Political Polarization in the United States: Predicting Party Identification from Social Attitudes and Demographics (1952-2024)

Introduction

The debate about whether American society is becoming more politically divided has been a major topic in political science for the last twenty years. Scholars mostly agree that elite polarization, which means the growing ideological differences between Democratic and Republican politicians, has increased a lot since the 1970s. However, there is still disagreement about how much this polarization is reflected in the general public's beliefs (Fiorina, Abrams, and Pope 2005; Abramowitz and Saunders 2008).

One important viewpoint, made by Fiorina (Fiorina, Abrams, and Pope 2005), argues that regular Americans are mostly moderate. According to this view, the appearance of polarization is caused by political leaders and the media influencing public opinion rather than real changes in people's attitudes. In contrast, Abramowitz (Abramowitz and Saunders 2008) believes that polarization is real and can be measured. He points out evidence that people who identify with a party have become more clearly aligned with certain ideological views. Mason (Mason 2018) adds that even if there haven't been big changes in people's policy preferences, the way people's partisan, ideological, and social identities align has led to a type of "social polarization." This social polarization has increased negative feelings between parties, visible not only in policy disagreements but also in how Democrats and Republicans feel about each other.

When looking at studies that separate opinion polarization (the differences in attitudes) from partisan sorting (how attitudes align with party labels), a more detailed picture forms. Baldassarri (Baldassarri and Gelman 2008) used data from the ANES (American National Election Studies) from 1972 to 2004. They found that while Americans did not become more extreme in their individual policy views, there was a significant increase in how much people's issue positions correlating with their party identification. They called this pattern "issue partisanship" rather than ideological constraint. Murphy (Murphy and Zhirnov 2020) continued this research through 2016 and confirmed that partisan alignment continued to increase, especially regarding cultural and moral issues.

Kim (Kim and Zilinsky 2022) took a different approach. They used logistic regression and machine learning methods to predict party identification based solely on demographic characteristics. Their finding revealed that demographics could only explain about 64% of partisan identifications, and that this predictive power has not increased over time. However, their analysis focused exclusively on demographics. The present study extends this framework by building a parallel attitudes-only track and tracking its accuracy across decades, a direct test of whether it is attitudes, not demographics, that are driving ideological sorting. Most recently, Ojer (Ojer et al. 2025) examined how different social groups have moved apart in ideological beliefs across demographic groups, revealing systematic divergence.

However, an important question remains regarding those who do not identify with either of the two main parties. The number of self-identified Independents has grown, and in 2024, more Americans identified as Independent than as Democrat or Republican (Gallup 2024). Nevertheless, research indicates that many of these so-called Independents are actually “undercover partisans” who lean strongly towards one party and vote accordingly (Klar and Krupnikov 2016). Klar and Krupnikov say that many Americans don’t like using partisan labels not because they have moderate views, but because they dislike political conflict. They find party politics divisive and want to stay away from it. This means that the unclear views of Independents might show that they are disengaged from politics instead of truly having a middle-ground ideology.

This study will look into both questions in two connected parts. The first part will track how the influences of attitudes and demographics on party identification have changed from the 1980s to the 2020s. This will be done using machine learning classification models applied to the ANES Cumulative Data File. The second part will focus on 2024 by asking who Democrats, Independents, and Republicans are today, and whether a genuine middle ground exists. To answer this question, the study will utilize multinomial classification models, misclassification analysis, and Latent Class Analysis without relying on party labels.

Research Questions and Hypotheses

This study is organized around two connected research questions.

Part 1: **How has the predictive power of social attitudes and socio-demographic characteristics for party identification among Americans changed between the 1980s and 2024, and what does this reveal about the process of ideological sorting?**

Part 2: **What do the people who don’t strongly identify with either major party look like, and is there really a true middle ground in the American voter base for 2024?**

These questions are operationalized through the following hypotheses:

- **H1:** Classification models trained on political attitudes will show higher accuracy in later periods than in earlier periods, consistent with increased ideological sorting.
- **H2:** Socio-demographic characteristics will become modestly more predictive of party identification over time, but will be consistently outperformed by attitude-based models.
- **H3:** Moral and cultural issues (abortion, homosexuality) will show the largest increase in predictive importance across decades, while economic issues will show smaller gains.
- **H4:** People who don’t strongly align with either political party will tend to be less interested in politics overall.
- **H5:** Despite identifying as Independent, most self-identified Independents will hold policy attitudes that align more closely with one of the two parties than with a genuine middle ground.

Data

American National Election Studies (ANES)

This study uses **the ANES Time Series Cumulative Data File (1948–2024)**, a large longitudinal survey that has tracked the political attitudes, behaviors, and demographic characteristics of American voters since 1948. The cumulative file combines 32 individual survey waves and includes 73,745 respondents across presidential and midterm election cycles.

The ANES is widely regarded as the most reliable dataset for studying American political behavior over time. Because the survey uses consistent question wording across decades, changes in our model results can be attributed to genuine shifts in how Americans think about politics, not to differences in how questions were asked.

Variables

The outcome variable is **party identification**, measured using the standard ANES 7-point scale ranging from Strong Democrat (1) to Strong Republican (7). For this analysis, we retain only respondents who identify as Democrats (1–3) or Republicans (5–7) and treat party identification as a binary outcome. Pure Independents (4) are excluded, following standard practice in the sorting literature (Baldassarri and Gelman 2008; Murphy and Zhirnov 2020). After excluding Independents and respondents with missing party identification, the analytical sample contains 63,659 respondents across 32 survey waves.

We use two sets of predictors. **The demographics track** includes age, gender, race, education, region, income, religion, church attendance, marital status, and employment status. **The attitudes track** includes 11 policy attitude items that have sufficient coverage across all decades from the 1980s onward: government job guarantee, aid to Black Americans, abortion, government spending, importance of religion, moral standards, traditional values, equal opportunity, and three items measuring attitudes toward inequality. These 11 items were selected because they are the only attitude variables asked consistently enough across survey waves to support decade-by-decade comparison, many attitude questions were introduced only in the 1980s or later, and variables with more than 60% missing values in any given decade were excluded.

Missing Data

A key challenge in longitudinal ANES analysis is that not all questions were asked in every survey wave. Missingness rates for some attitude variables exceed 50% in earlier decades. To address this, we use Multiple Imputation by Chained Equations (MICE) with predictive mean matching, generating five imputed datasets per decade. Imputation is conducted separately within each decade to preserve historical context and prevent information from later periods from influencing earlier ones. This approach is preferable to simply dropping observations with missing values, which would disproportionately remove respondents from earlier survey waves and bias the sample toward more recent years.

Methods

This study uses a decade-by-decade classification framework to track how well attitudes and demographics predict party identification from the 1980s through the 2020s. For each decade, we train separate models on 70% of the data and evaluate them on the remaining 30%. The key insight behind this approach is straightforward: if ideological sorting has increased over time, attitude-based models should become more accurate in later decades. This follows Kim and Zilinsky (2022), who argued that predictive accuracy is a more direct and intuitive measure of sorting than regression coefficients alone.

Two analytical tracks are maintained separately throughout. The demographics track predicts party identification from socio-demographic characteristics only. The attitudes track predicts party identification from policy attitude items only, with MICE imputation for missing values. Comparing the two tracks allows us to assess whether it is attitudes or demographics that are driving the sorting trend.

Models

We apply three classification algorithms to each track. **Logistic regression serves as the interpretable baseline.** Its coefficients tell us directly how each predictor relates to Republican versus Democratic identification, and we use these coefficients to construct heatmaps that track how the importance of specific attitudes has shifted across decades. **Random forests serve as the primary performance benchmark**. By averaging over many decision trees, random forests reduce overfitting and typically outperform single models. We use 500 trees throughout, with hyperparameter tuning where appropriate. **Decision trees provide a visual summary of the most important classification rules**, showing, for example, which combinations of race, religion, and education most clearly predict party identification in a given decade.

We also apply LASSO logistic regression to the attitudes track as a variable selection tool. LASSO penalizes the size of coefficients and can shrink less important predictors to zero, helping us identify which attitude variables carry the most independent predictive signal across decades.

Exploratory Data Analysis

Before turning to the machine learning models, we first examine descriptive trends in party identification and demographic sorting. These patterns provide important context for understanding what the models are trying to predict, and why we might expect their accuracy to change over time.

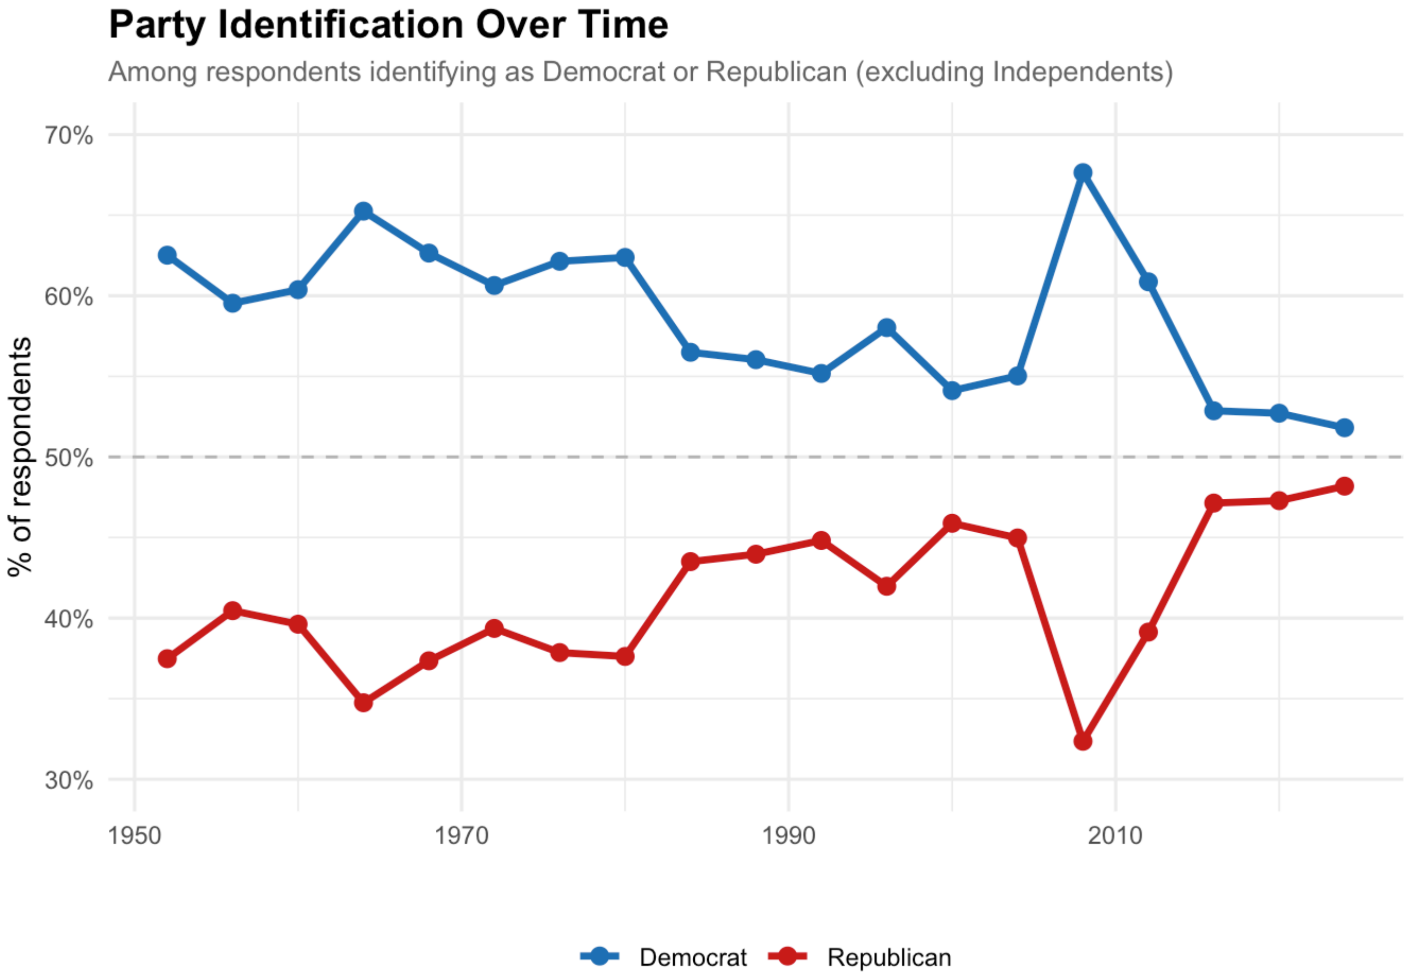

The overall trend in party identification shows a substantial Democratic advantage through the 1950s–1970s, reflecting the post-New Deal coalition that dominated American politics for much of the twentieth century. The gap has narrowed considerably since the 1980s, with parties approaching near-parity in the 2020s, consistent with the nationalization of American politics documented by Abramowitz and Saunders (2008), as national partisan identity has grown stronger, the old regional and class-based coalitions that sustained Democratic dominance have gradually eroded.

How have different demographic groups changed over time? We look at party affiliation based on six key factors: **education, gender, income, religion, church attendance, region, and marital status**. These patterns reveal which demographic divides have become stronger or weaker in partisan politics.

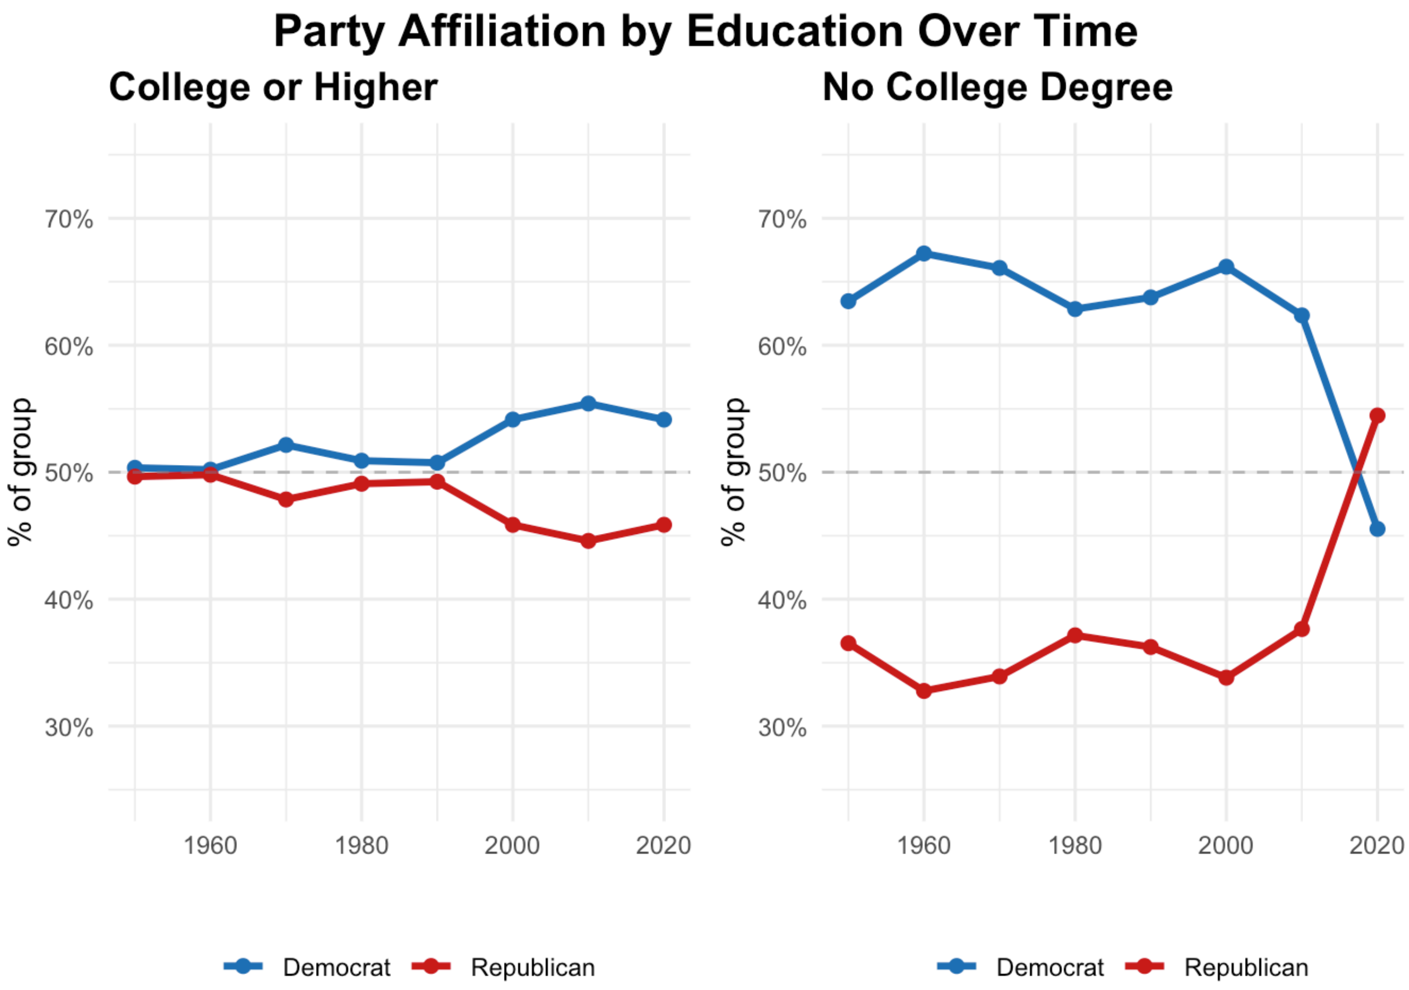

The education reversal is a significant change in American politics. College-educated voters, who used to mostly support Republicans, have increasingly leaned toward Democrats since the 1990s. In contrast, non-college-educated voters have shifted the other way. This growing divide based on education has been widely studied (Abramowitz and Saunders 2008; Ojer et al. 2025) and became especially noticeable between 2010 and 2020.

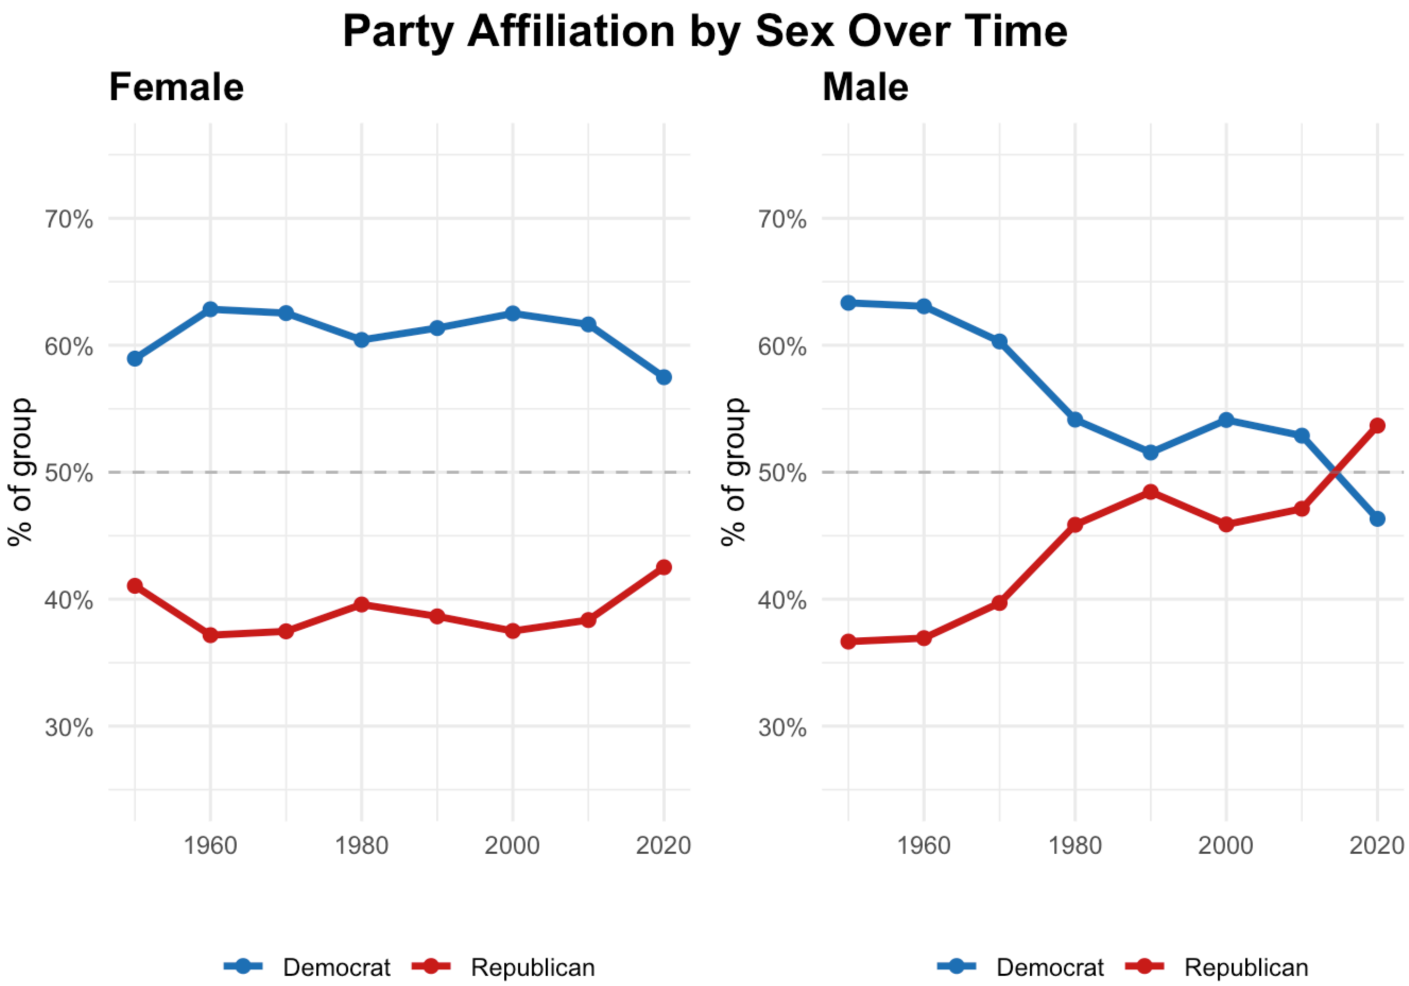

The gender gap in partisanship, while modest in the 1950s–1960s, has grown over time. Women have consistently leaned more Democratic than men since the 1980s, a pattern linked to diverging views on social issues, reproductive rights, and the role of government (Mason 2018).

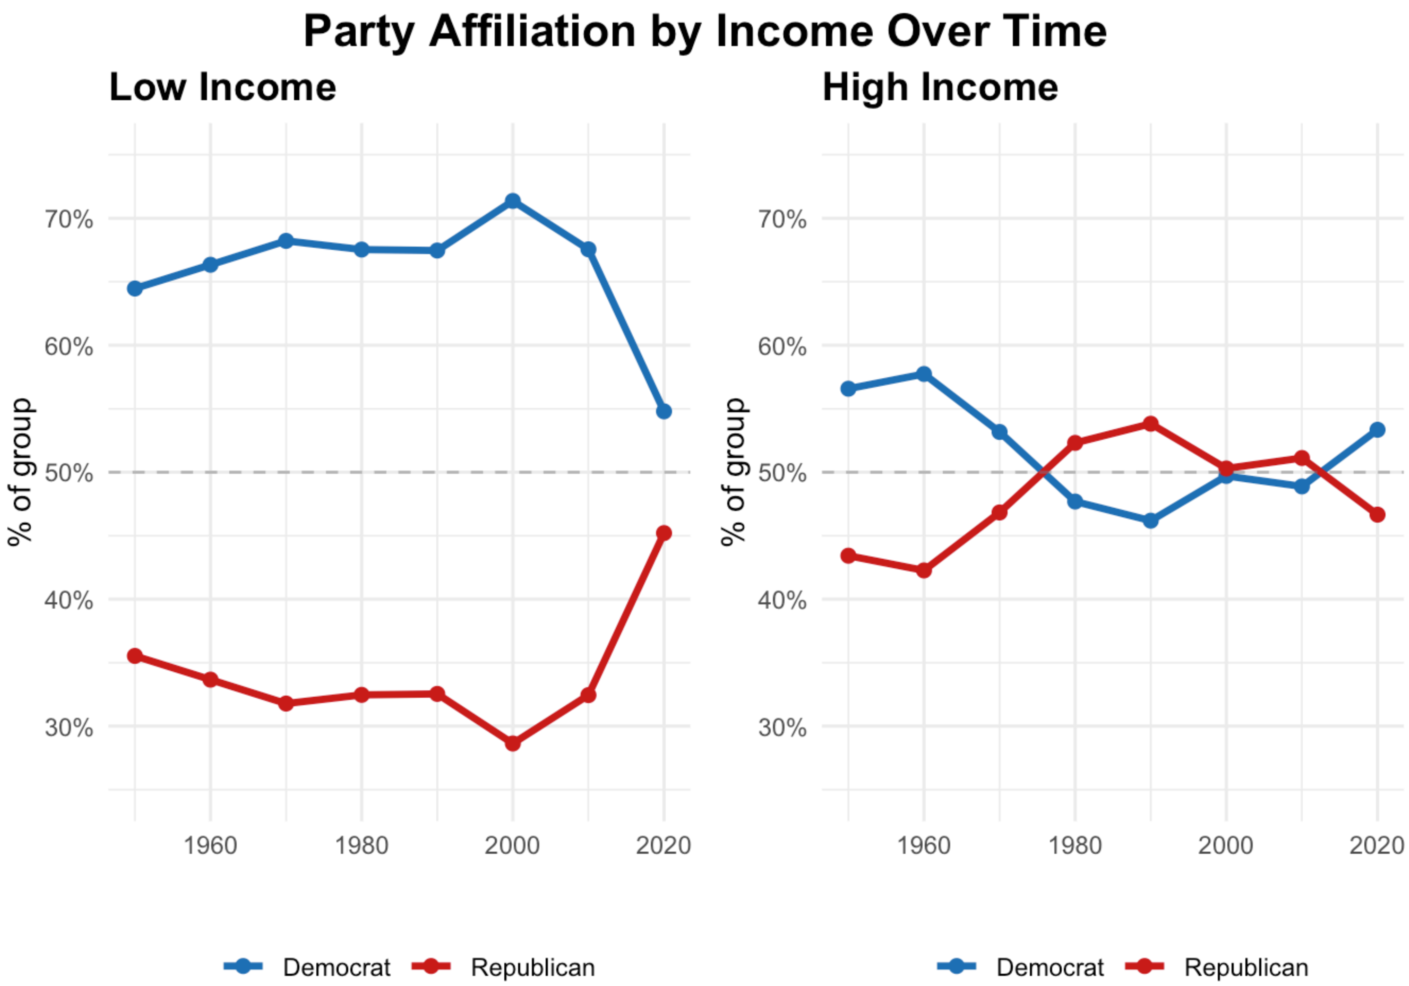

Low-income voters have typically supported Democrats, but since 2000, that support has decreased and nearly equalized by the 2020s. Conversely, high-income voters have shifted from a strong Republican preference to a more balanced support between the parties. Overall, income is becoming less important in predicting political party affiliation.

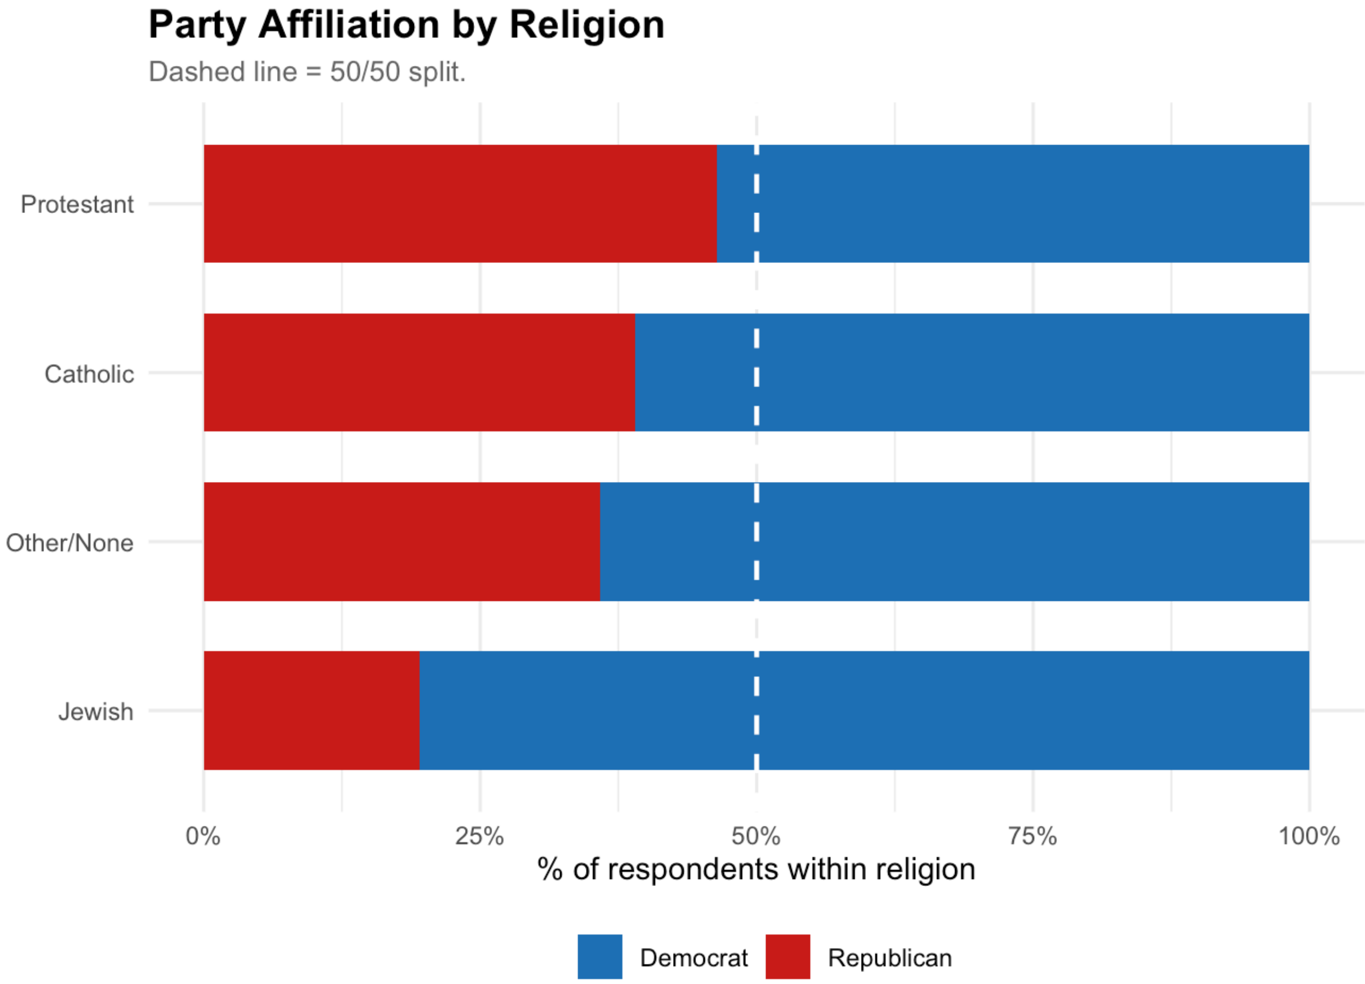
Party Affiliation by Religion
Dashed line = 50/50 split.
Protestant
Catholic
Other/None
Jewish
0%
25%
50%
75%
100%
% of respondents within religion
Democrat
Republican

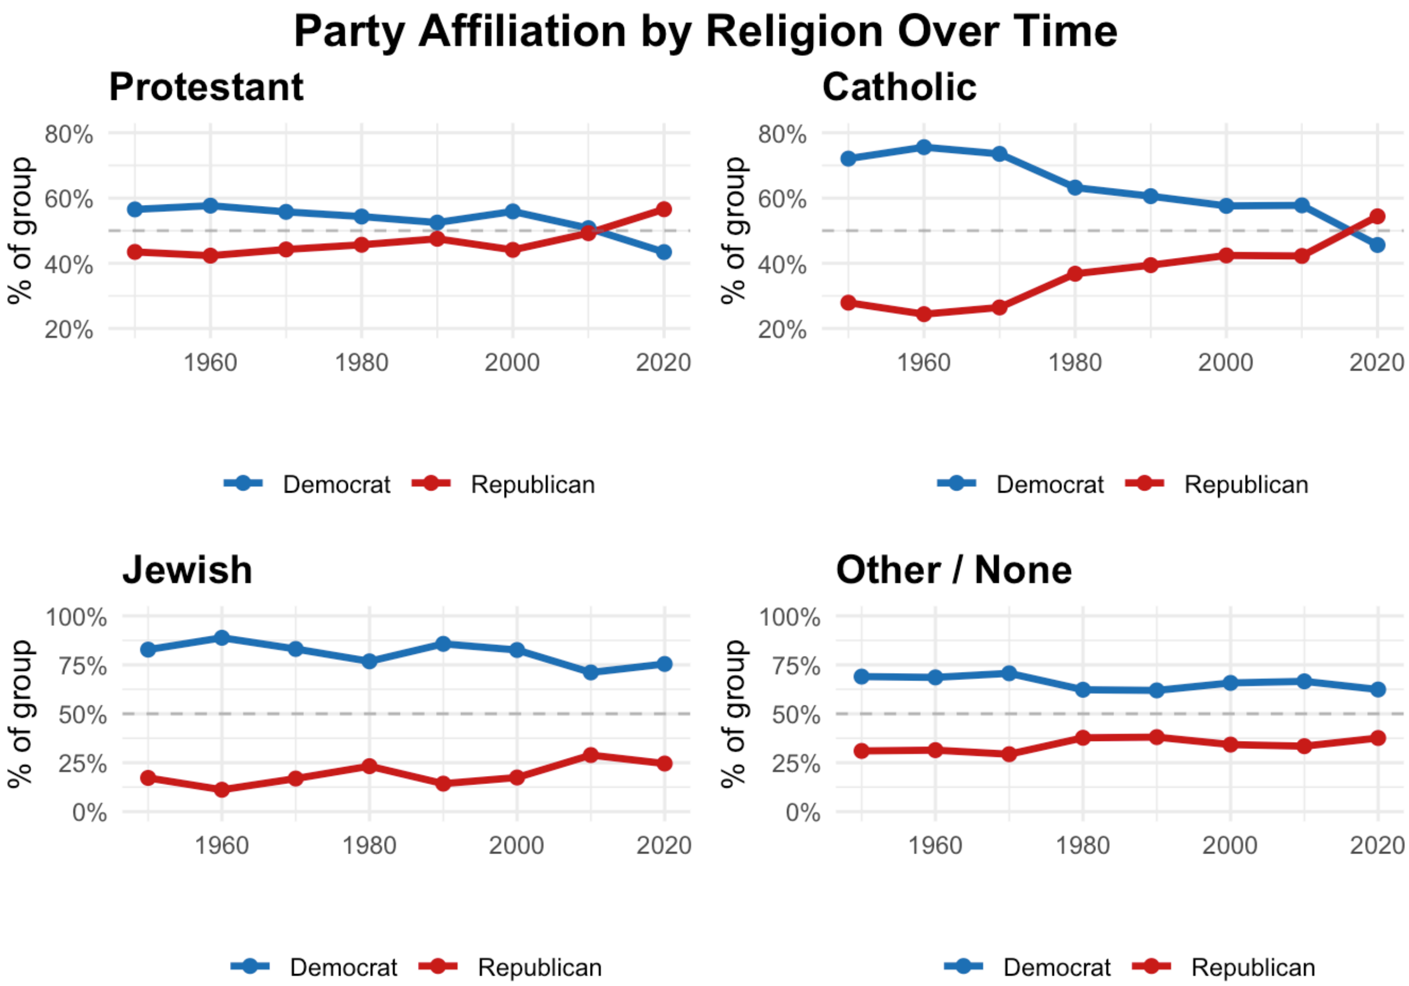

Religion is a key factor in how people vote. Many white Protestants and more Catholics are leaning Republican, reflecting the “culture war” sorting described by Baldassarri and Gelman (2008) and Murphy and Zhirnov (2020), while Jewish voters and those without religious ties mostly support Democrats. The number of religiously unaffiliated voters, often called “nones,” has increased, leading to a growing divide between religious and non-religious voters since the 1990s.

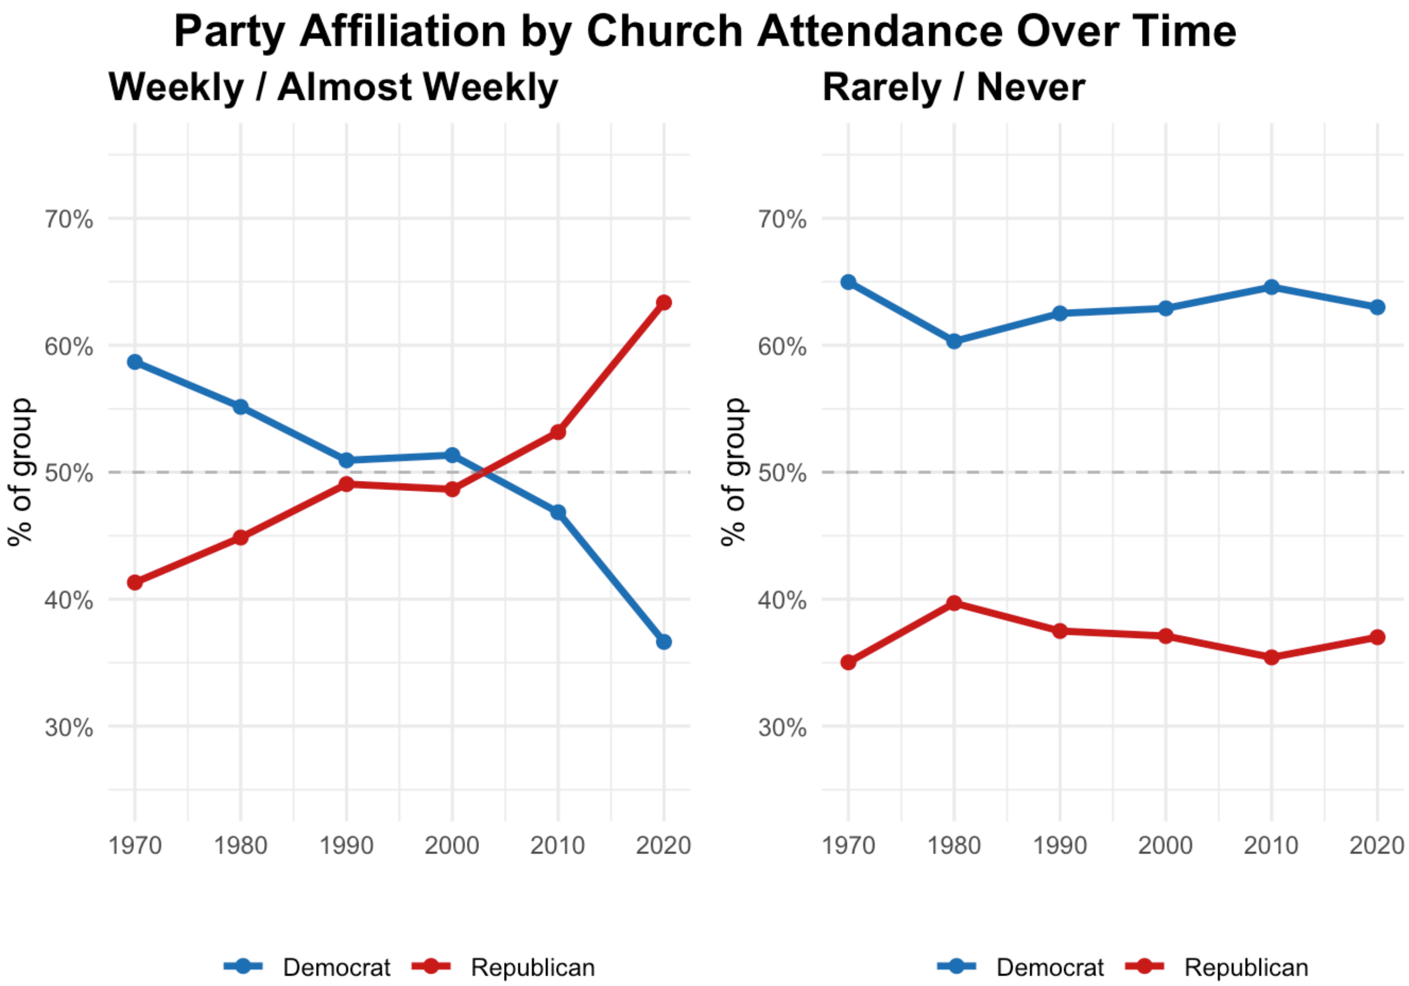

Church attendance has become a strong indicator of whether someone identifies as a Republican. Since the 1970s, people who attend church weekly have increasingly leaned toward the GOP, especially in the 1990s with the rise of the Christian right. In contrast, those who rarely or never attend church have shifted away from the Republican Party, widening the gap based on religious attendance.

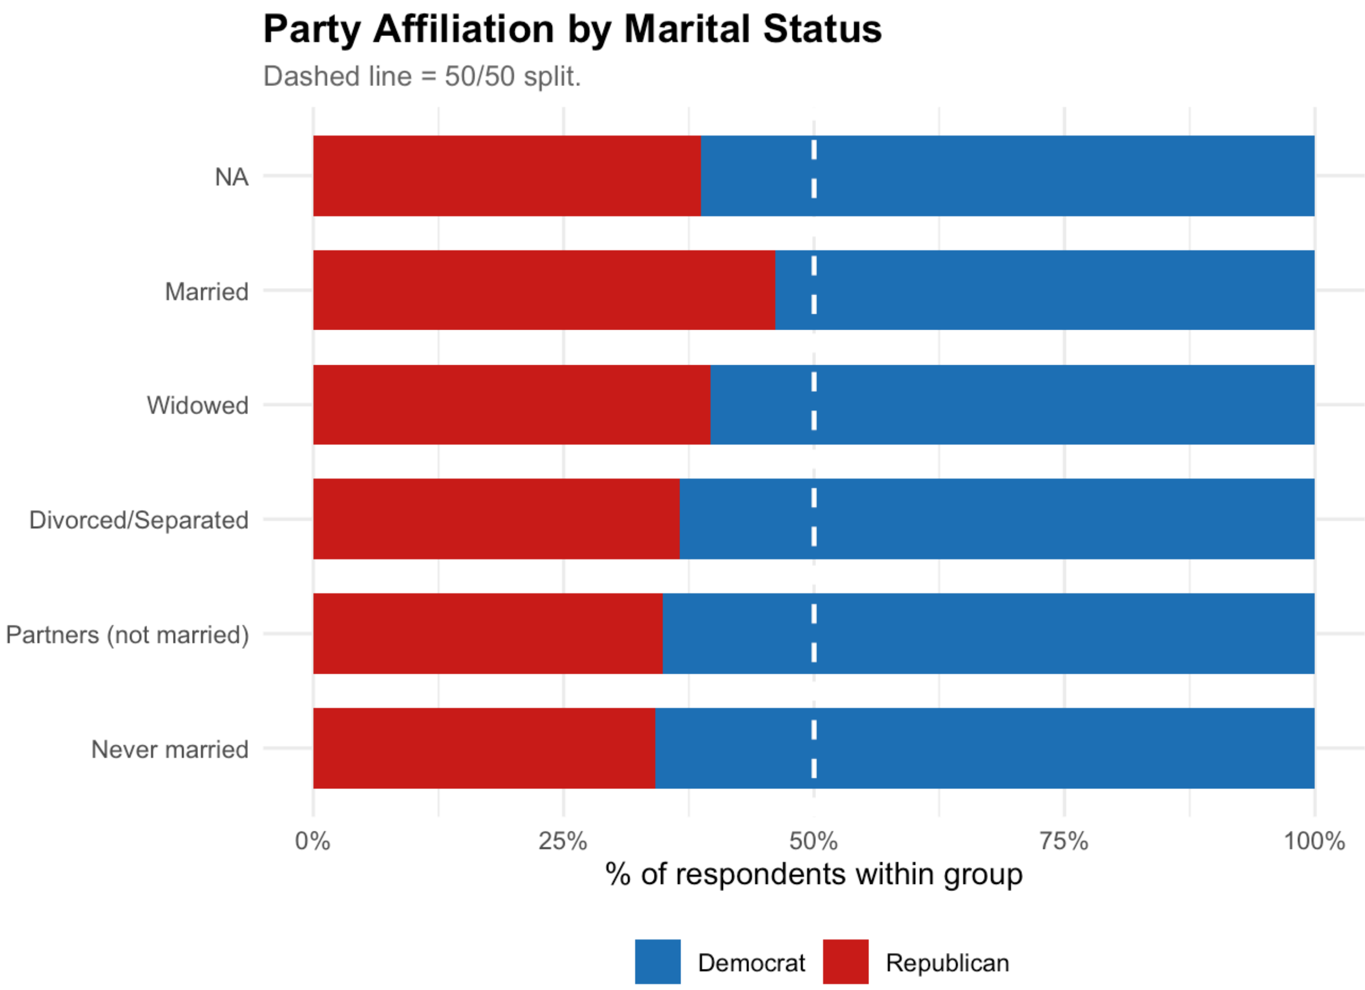

Marital status shows some differences in political views. Married people tend to favor Republicans a bit, while those who are never married or living together lean more towards Democrats. However, these differences aren’t very significant, and marital status is not a major factor in predicting political preferences.

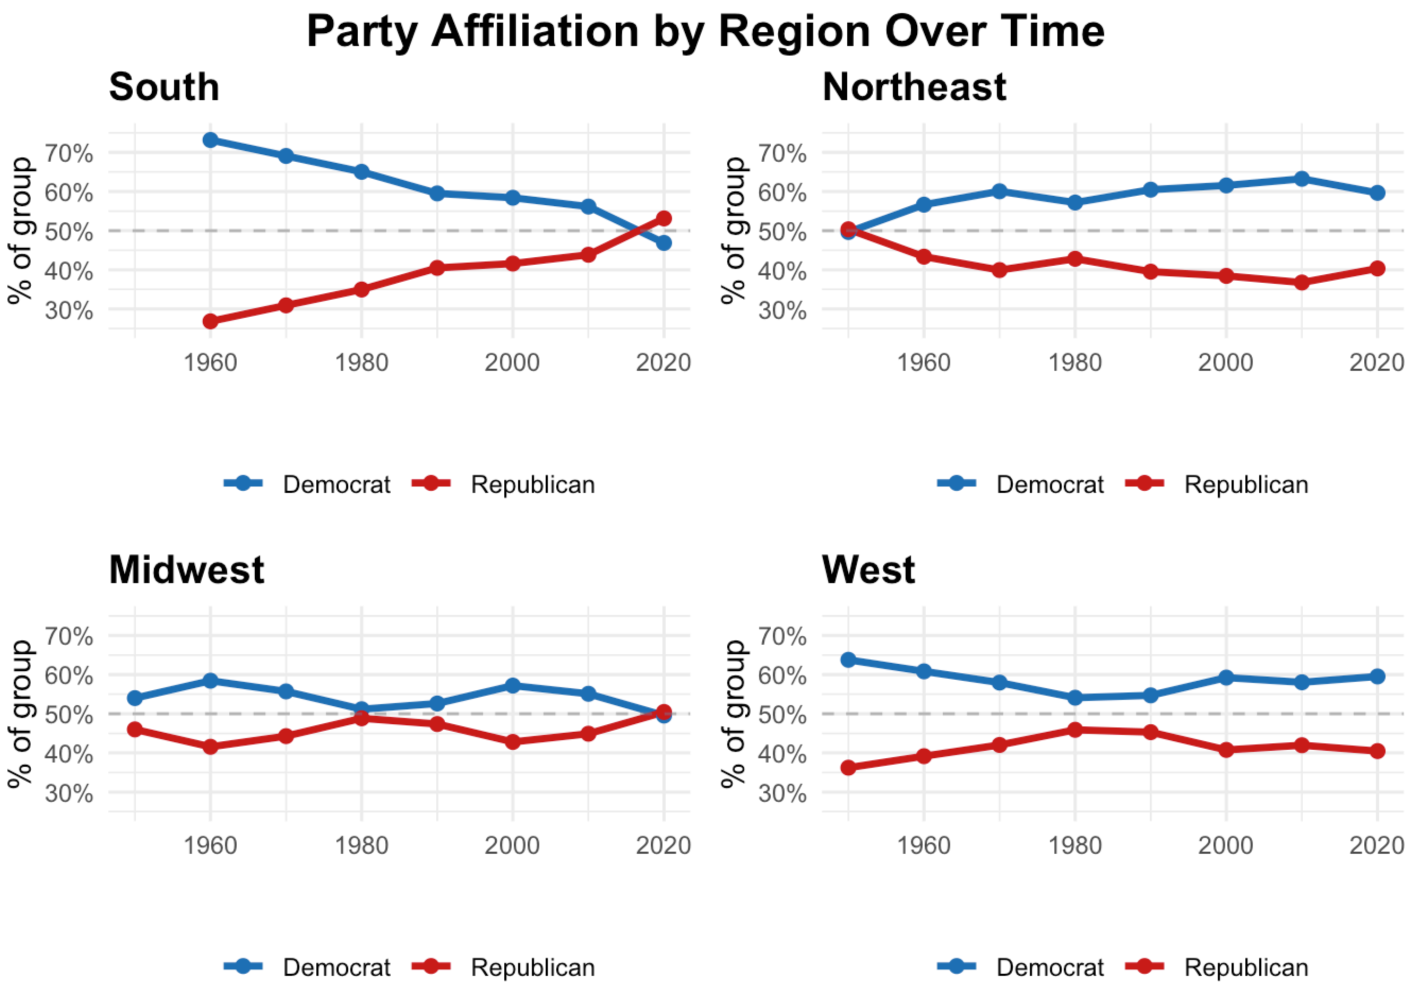

Over the past 70 years, there has been a big shift in how different regions in the U.S. vote. The South used to be mostly Democratic but has become strongly Republican since the 1990s, largely because of the Civil Rights Act. In contrast, the Northeast has changed from competitive to mainly Democratic. This shows that national politics are becoming more important than local voting habits.

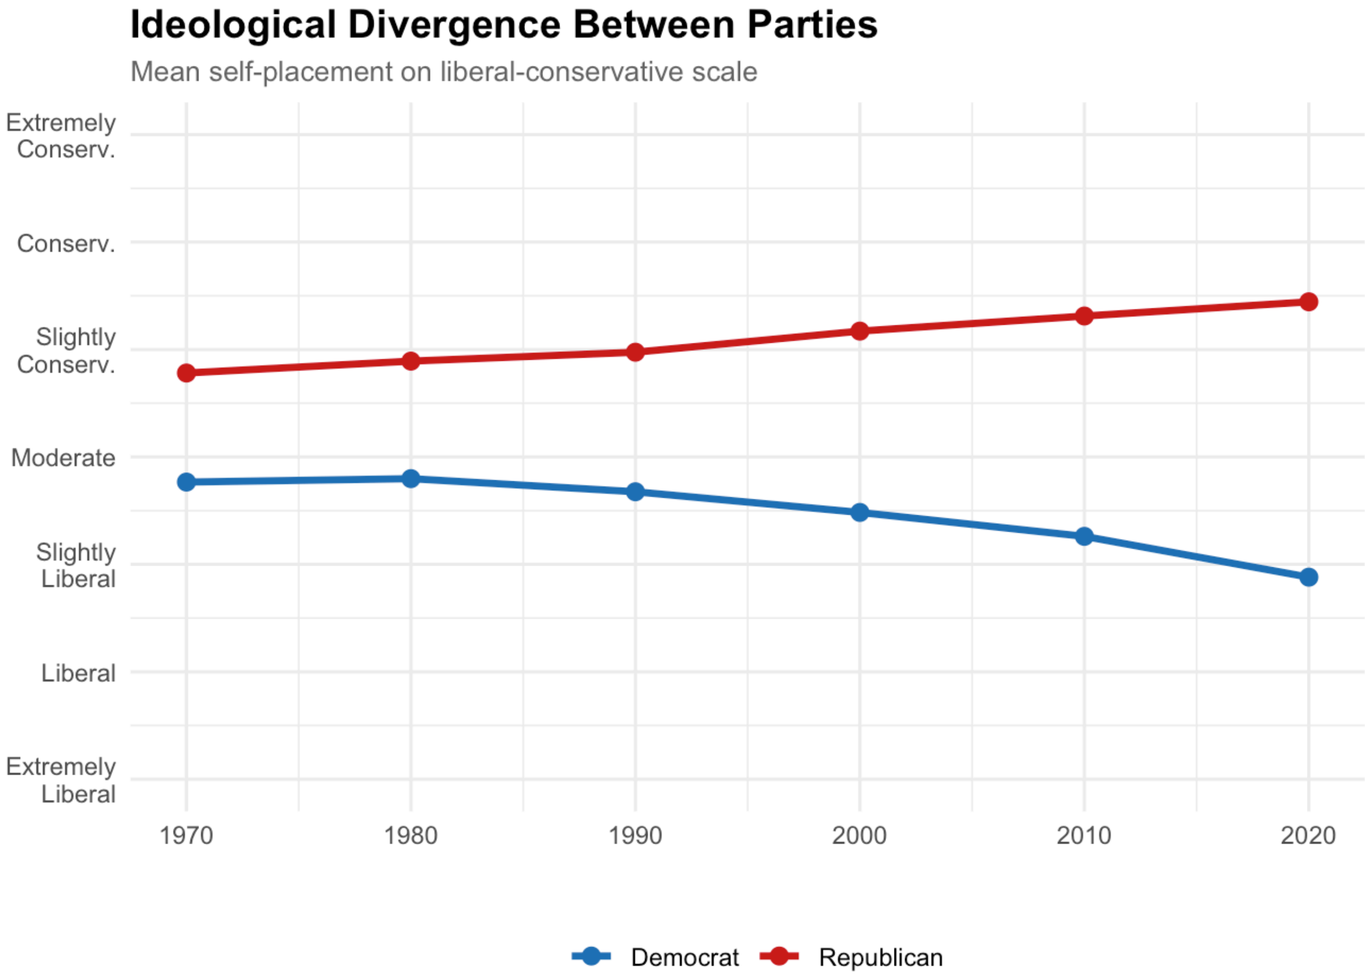

The ideological divergence chart shows that political polarization has increased over time. In the 1970s, Democrats identified as “Moderate” and Republicans as “Slightly Conservative,” indicating a small gap between the two. By the 2020s, Democrats have moved to “Slightly Liberal,” while Republicans have become more “Conservative.” This trend supports the Abramowitz and Saunders (2008)’s idea of partisan sorting, meaning people are aligning more closely with their party’s values. It’s important to note that this chart reflects self-placement on a single dimension; the machine learning models below use issue-specific attitudes to assess *actual* sorting behavior, not just self-identification.

Beyond demographic sorting, scholars have increasingly drawn attention to a parallel phenomenon: the rise of affective polarization, or growing emotional hostility between partisan groups. Unlike ideological polarization, which concerns what people believe, affective polarization concerns how people feel about those on the other side (Iyengar et al. 2019).

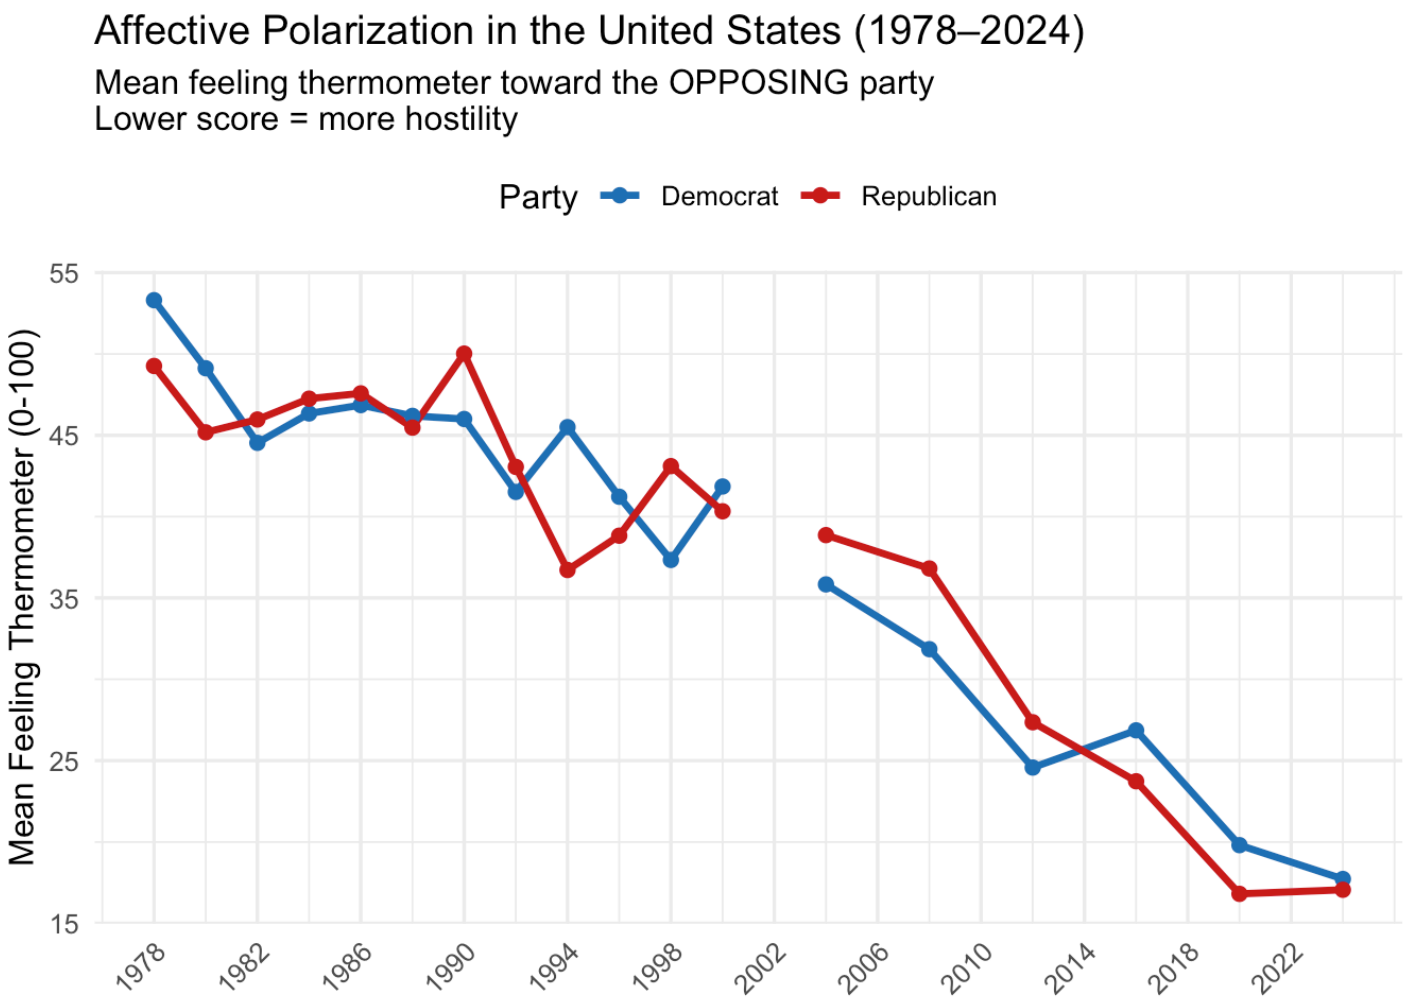

Research shows that feelings toward the opposing political party have declined sharply since the late 1970s. Ratings that Democrats give to Republicans and those Republicans give to Democrats have both dropped steadily for the past 40 years, reaching very low levels by the 2020s. This trend aligns with what Iyengar et al. (2019) describes as affective polarization: Americans are not necessarily more extreme in their policy views, but they have become more hostile toward the other party. The decline sped up after the mid-1990s, around the time of partisan shifts in religion and education. There is no indication that this trend will change by 2024. The emotional gap between the two parties has widened almost perfectly, suggesting that as party identities have become stronger, feelings of distance between the parties have also grown. There is no data for 2002 due to a lack of thermometer items in that survey.

Demographics Track

We begin with the demographics track, which uses only socio-demographic characteristics — age, gender, race, education, religion, income, region, and church attendance, to predict party identification. This track tests H2: **whether demographic characteristics have become more predictive of party identification over time**. The results provide a useful baseline against which to compare the attitude-based models in the following section.

Logistic Regression

The logistic regression model demonstrates a steady improvement in predicting accuracy over the decades, increasing from 63.6% in the 1980s to 71.4% in the 2020s. AUC scores also show an upward trend, rising from 71.3 to 77.6 during the same time. This suggests that demographic factors are becoming more linked to party identification, which aligns with H2. However, the overall accuracy is still limited. Even in the 2020s, about one in three people are misclassified based on demographics alone. Knowing someone's age, race, income, or education doesn't provide a clear picture of their political party affiliation, supporting the findings of Kim and Zilinsky (2022), who indicated that demographics reveal surprisingly little about how people vote.

Logistic Regression Performance by Decade (Demographics Track)

decade	accuracy	f1	auc	n_train	n_test
1980s	63.6	61.5	71.3	3258	1395
1990s	66.6	59.6	72.5	4830	2069
2000s	70.3	59.7	75.2	2774	1187
2010s	71.3	66.0	78.4	5574	2388
2020s	71.4	71.8	77.6	6241	2673

Decision Tree

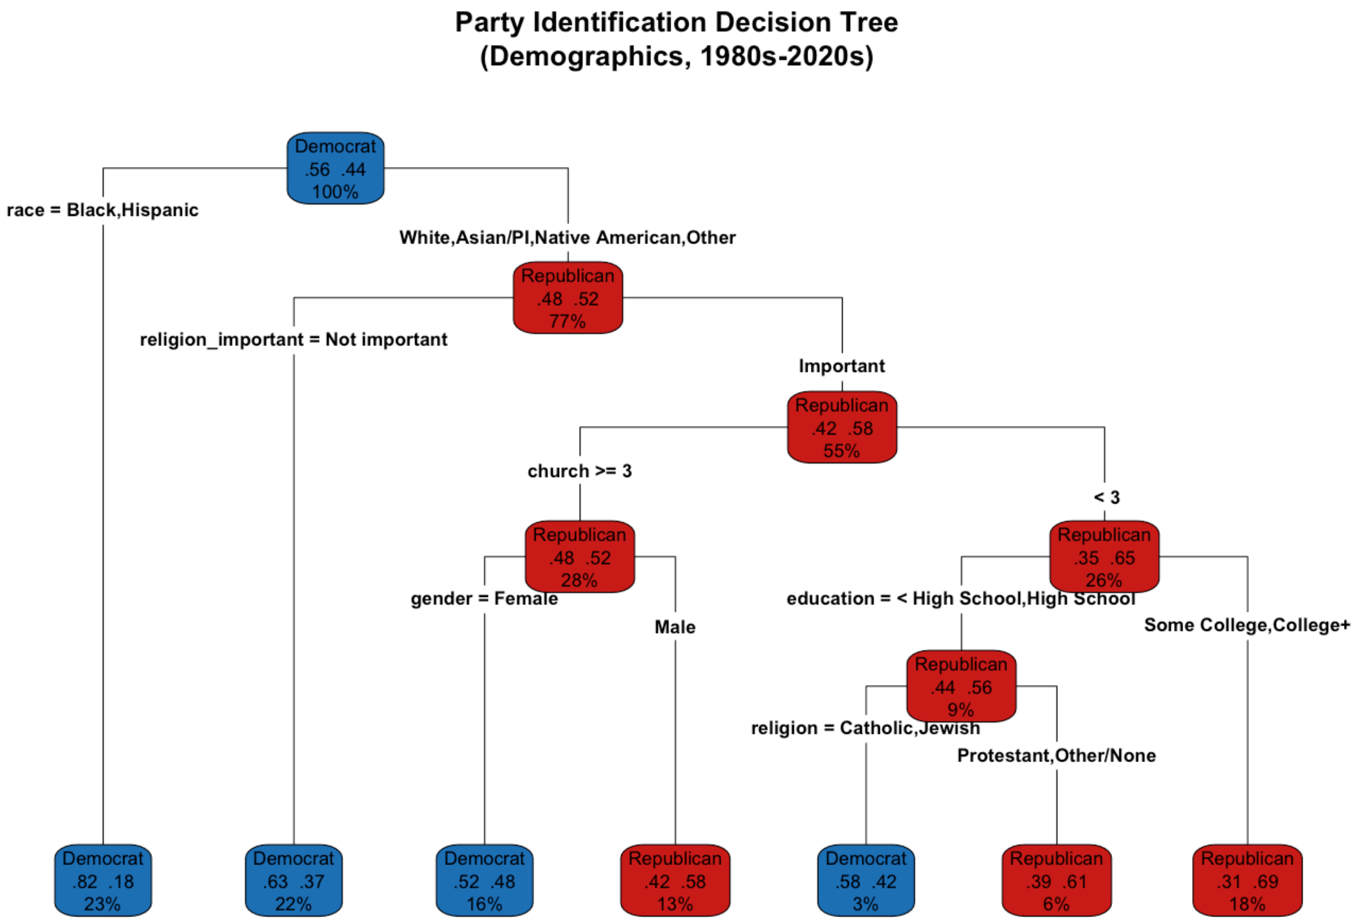

The decision tree shows how different demographic groups tend to vote. It highlights that Black and Hispanic respondents are mostly classified as Democrats, while white respondents are further divided based on factors like religious beliefs, church attendance, gender, and education. This pattern reflects the way race influences American politics and the growing role of religious identity.

However, decision trees have a drawback: they can be unstable. Small changes in the data or how it's split into training and testing sets can lead to different results throughout the tree. This is why we will look at Random Forest next, which helps address this issue.

Random Forest

The random forest model supports the findings from the other methods. The accuracy increased from 63.9% in the 1980s to 71.4% in the 2020s, with a small drop to 68.4% in the 2000s, likely due to a smaller sample size from that decade (n_train = 2,774). The F1 score also improved significantly, going from 59.7 in the 1980s to 71.2 in the 2020s. This indicates that the model has become better at identifying both Democrats and Republicans, not just the majority party.

Random Forest Performance by Decade (Demographics Track)

decade	accuracy	f1	auc	n_train	n_test
1980s	63.9	59.7	69.9	3258	1395

1990s	69.8	62.0	75.3	4830	2069
2000s	68.4	54.9	74.9	2774	1187
2010s	70.6	64.8	78.2	5574	2388
2020s	71.4	71.2	77.6	6241	2673

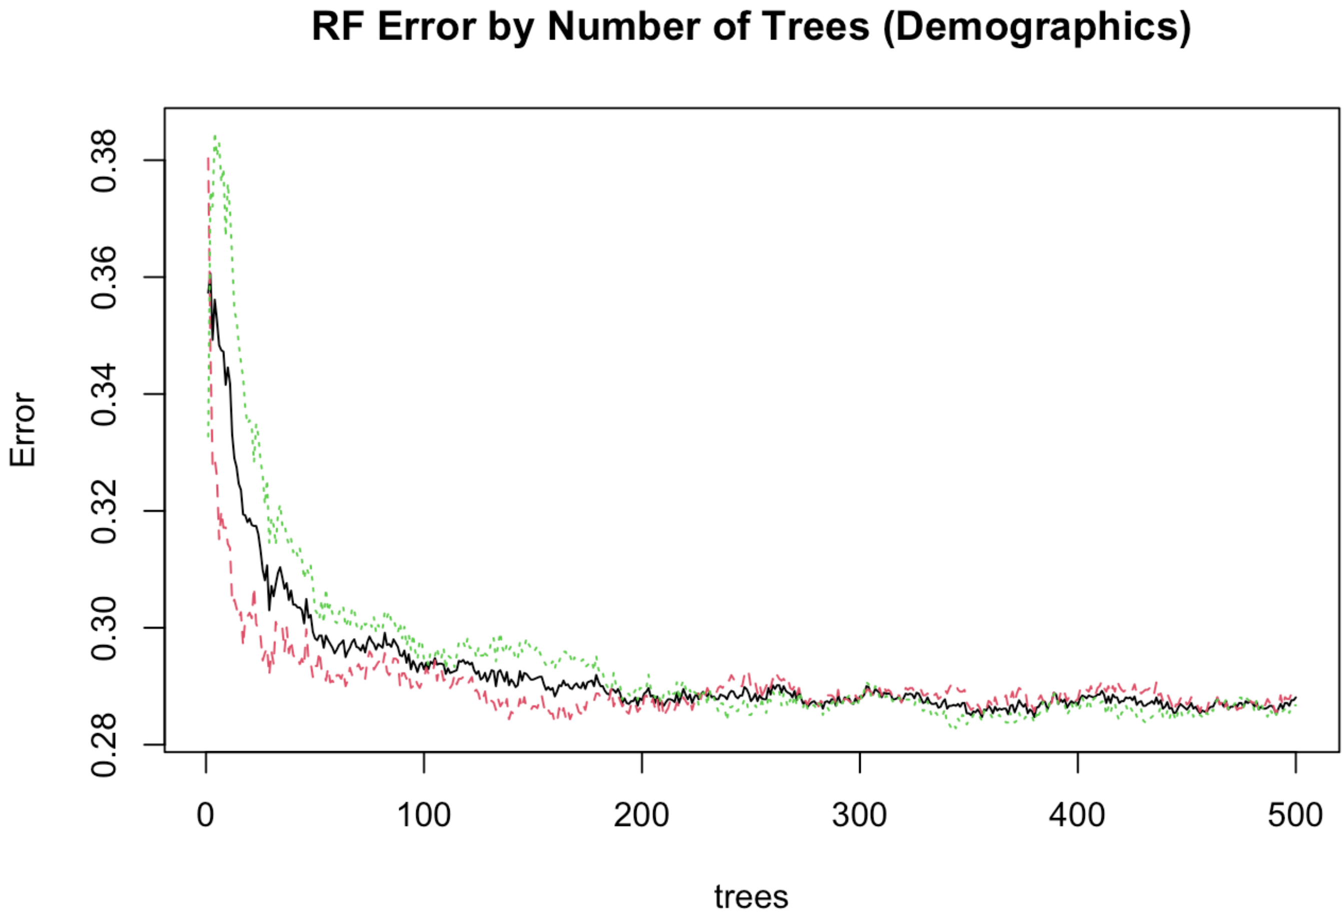

The plot above shows the out-of-bag error rate as a function of the number of trees, using the 2020s decade as a representative example. The error curve flattens well before 500 trees, confirming that ntree = 500 is a sufficient and stable choice. Increasing the number of trees further would not meaningfully improve model performance.

Demographics Summary Plot

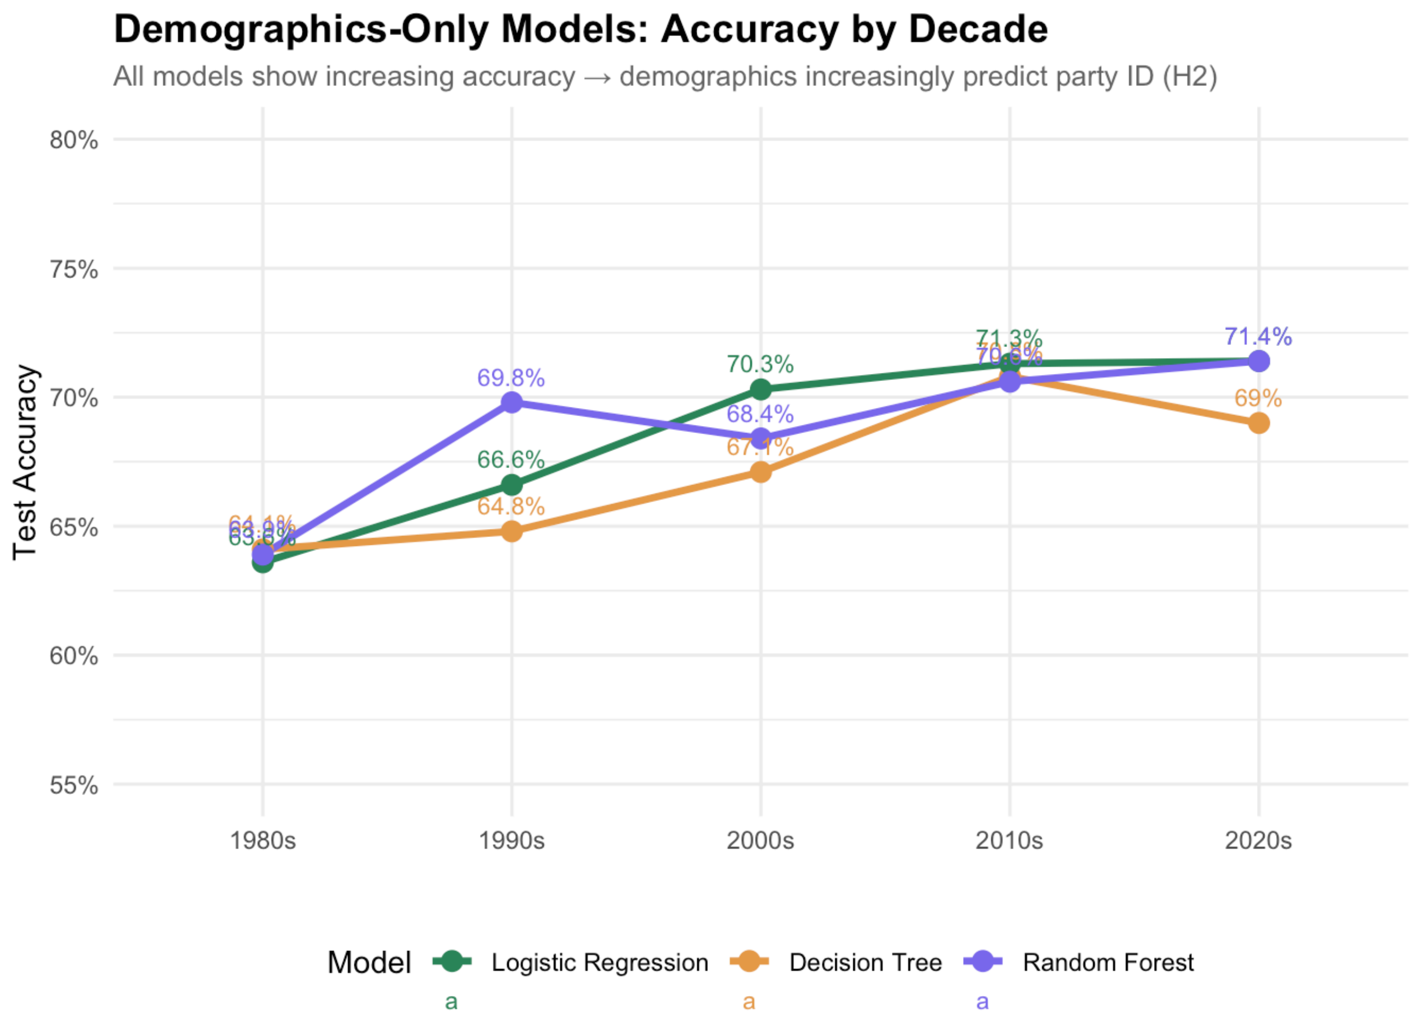

Demographics-only accuracy increases from approximately 64% in the 1980s to 71% in the 2020s across all three model types. This ~7 percentage point gain over four decades suggests that demographic characteristics have become modestly more predictive of party identification, consistent with H2. However, the rate of increase is smaller than for the attitudes track (below), suggesting that attitudinal sorting, rather than demographic realignment, is the primary driver of the overall polarization trend.

Attitudes Track

We now focus on how people's policy attitudes influence their party identification instead of looking at demographics. This is the main part of our analysis. If ideological sorting has grown over time, then understanding someone's beliefs about issues like abortion, government spending, or racial equality should better predict their political party in 2024 compared to the 1980s.

Logistic Regression with MICE Imputation

The results are significant. Accuracy has increased from 67.8% in the 1980s to 85.9% in the 2020s—an 18 percentage point increase over four decades. The AUC scores reflect a similar trend, rising from 73.7 to 93.0. In comparison, demographic tracking has improved by only about 7 percentage points during the same period. The attitude-based models are not only better but are also improving at more than twice the rate. By the 2020s, the F1 score reaches 85.1, indicating that the model reliably classifies both Democrats and Republicans with high and nearly equal accuracy. This provides strong support for the hypothesis that ideological sorting has increased significantly, driven by attitudes rather than demographics.

decade	model	accuracy	f1	auc	n_total
1980s	Logistic Regression	67.8	58.8	73.7	8183
1990s	Logistic Regression	69.7	62.0	75.9	8146
2000s	Logistic Regression	75.4	68.0	82.6	6032
2010s	Logistic Regression	82.0	78.3	89.3	8766
2020s	Logistic Regression	85.9	85.1	93.0	12380

Logistic Regression Coefficient Heatmap

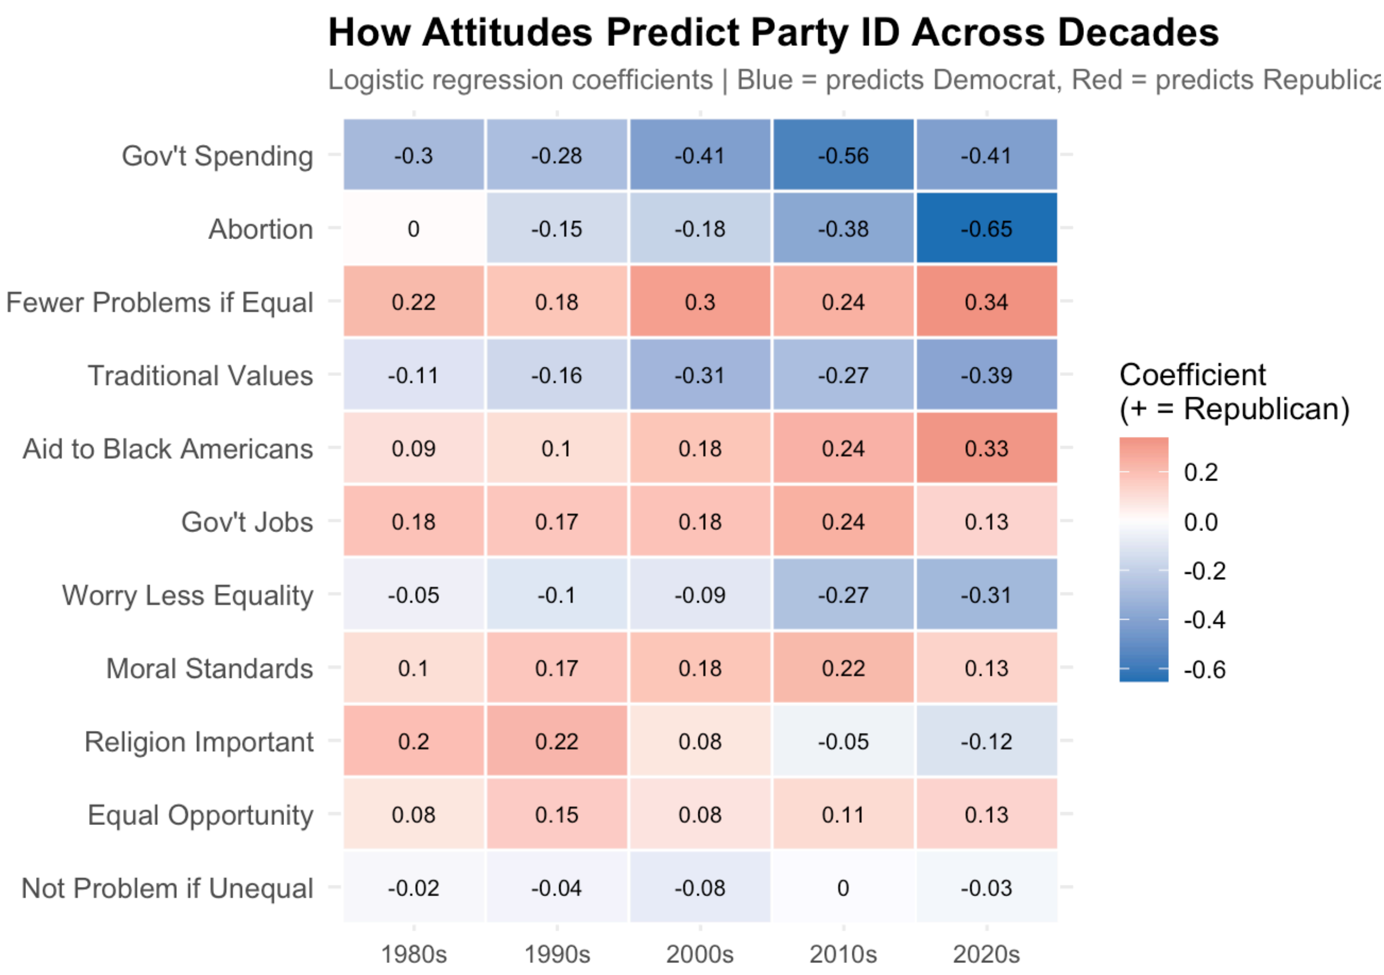

To understand what factors are leading to improved Democratic identification, we look at how individual attitudes have changed over the decades. The heatmap shows some key trends. **Abortion** has seen the biggest shift, with its impact growing from almost zero in the 1980s to a strong negative coefficient of

-0.65 in the 2020s. This makes it the strongest predictor of Democratic identification by the end of this period (a negative value means more pro-choice views are linked to more Democratic identification). **Government spending preferences** have also become more supportive of Democratic views.

Meanwhile, attitudes around **traditional values** and **moral standards** have consistently leaned Republican, with their influence increasing over time. In contrast, economic factors such as government job guarantees and equal opportunity have remained relatively stable and modest in impact across all decades. This partially supports the idea that cultural and moral issues have become more important in shaping political identification than economic issues, but this trend isn't consistent for all factors.

LASSO

The LASSO results support the findings from the logistic regression. Unlike older models that would reduce weak predictors to zero, LASSO keeps all 11 variables in most decades. This shows that each attitude item has its own predictive value. The agreement between these two different methods boosts our confidence that the trend we see is a real change in the relationship between attitudes and political parties, not just due to how the model is set up.

decade	accuracy
1980s	67.6
1990s	69.1
2000s	77.7
2010s	83.0
2020s	85.6

Random Forest with MICE

decade	model	accuracy	f1	auc	n_total
1980s	Random Forest	68.3	58.6	73.4	8183
1990s	Random Forest	69.1	60.6	75.1	8146
2000s	Random Forest	75.3	67.5	81.7	6032
2010s	Random Forest	81.7	77.6	89.1	8766
2020s	Random Forest	85.8	84.9	93.0	12380

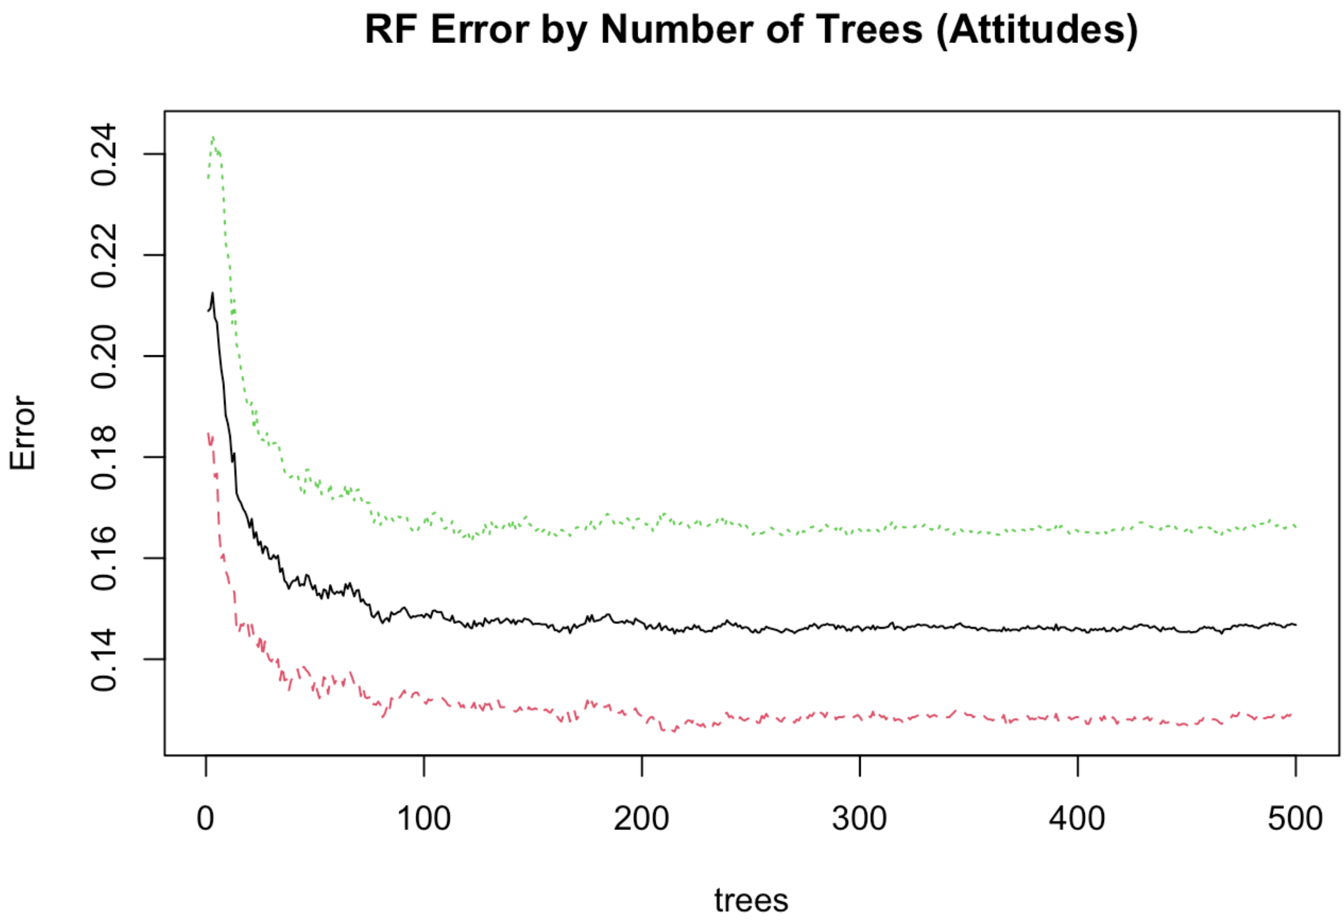

The plot above shows the out-of-bag error rate as a function of the number of trees, using the 2020s decade as a representative example. The error curve flattens well before 500 trees, confirming that ntree = 500 is a sufficient and stable choice. Increasing the number of trees further would not meaningfully improve model performance.

Model Accuracy Comparison: Attitudes vs. Demographics Track (%)

Model Accuracy by Decade, Track, and Algorithm (%)

Decade	Track	Model	Accuracy
1980s	Attitudes	Logistic Regression	68.9
1980s	Attitudes	Random Forest	69.6
1980s	Demographics	Logistic Regression	63.6
1980s	Demographics	Random Forest	64.2
2020s	Attitudes	Logistic Regression	85.5

2020s	Attitudes	Random Forest	85.0
2020s	Demographics	Logistic Regression	70.5
2020s	Demographics	Random Forest	70.3

In the 1980s, all four models performed modestly, showing that party identification wasn't closely linked to attitudes or demographics at the time. The Attitudes models slightly outperformed the Demographics models, with Logistic Regression achieving 68.9% accuracy and Random Forest 69.6%, while the Demographics models had 63.6% and 64.2% accuracy. All models performed significantly better than random guesses, indicating they identified real patterns rather than just noise.

Interestingly, both Attitudes models had more difficulty correctly identifying Republicans compared to Democrats. They accurately classified about 550–540 Republicans but misclassified around 488–498 as Democrats. This pattern suggests that in the 1980s, Republican identity wasn't as strongly tied to policy attitudes, reflecting an early stage of ideological sorting. The Demographics models showed a similar trend, with moderate but overall weak classification for both parties.

By the 2020s, the difference between the political landscape and that of the 1980s is clear. The models based on attitudes show significant improvement, achieving about 85% accuracy with Logistic Regression and Random Forest. They accurately identify Democrats and Republicans, indicating that people's policy beliefs are now strong indicators of their party affiliation.

Models based on demographics have also improved, reaching around 70% accuracy, but the gap between the two has widened. In the 1980s, attitude-based models were only slightly more accurate than demographic ones by about 5 points, but by the 2020s, this difference has grown to about 15 points.

Attitudes Summary Plot

The attitudes track has shown a significant increase in predictive accuracy, rising from about 68% in the 1980s to 86% in the 2020s—an 18 percentage point improvement. In comparison, the demographics track only improved by about 7 percentage points. This suggests that ideological beliefs, rather than demographic factors, are the main drivers of partisan sorting. Both logistic regression and random forest models show similar trends, confirming the robustness of these findings.

This aligns with research by Murphy and Zhirnov (2020), which noted increasing issue partisanship from 1972 to 2016, and Ojer et al. (2025), who found that different demographic groups are diverging ideologically. The accuracy of 86% in the 2020s greatly surpasses the 64% limit found in demographics-only models by Kim and Zilinsky (2022), reinforcing the idea that attitudes are what shape modern partisan identity.

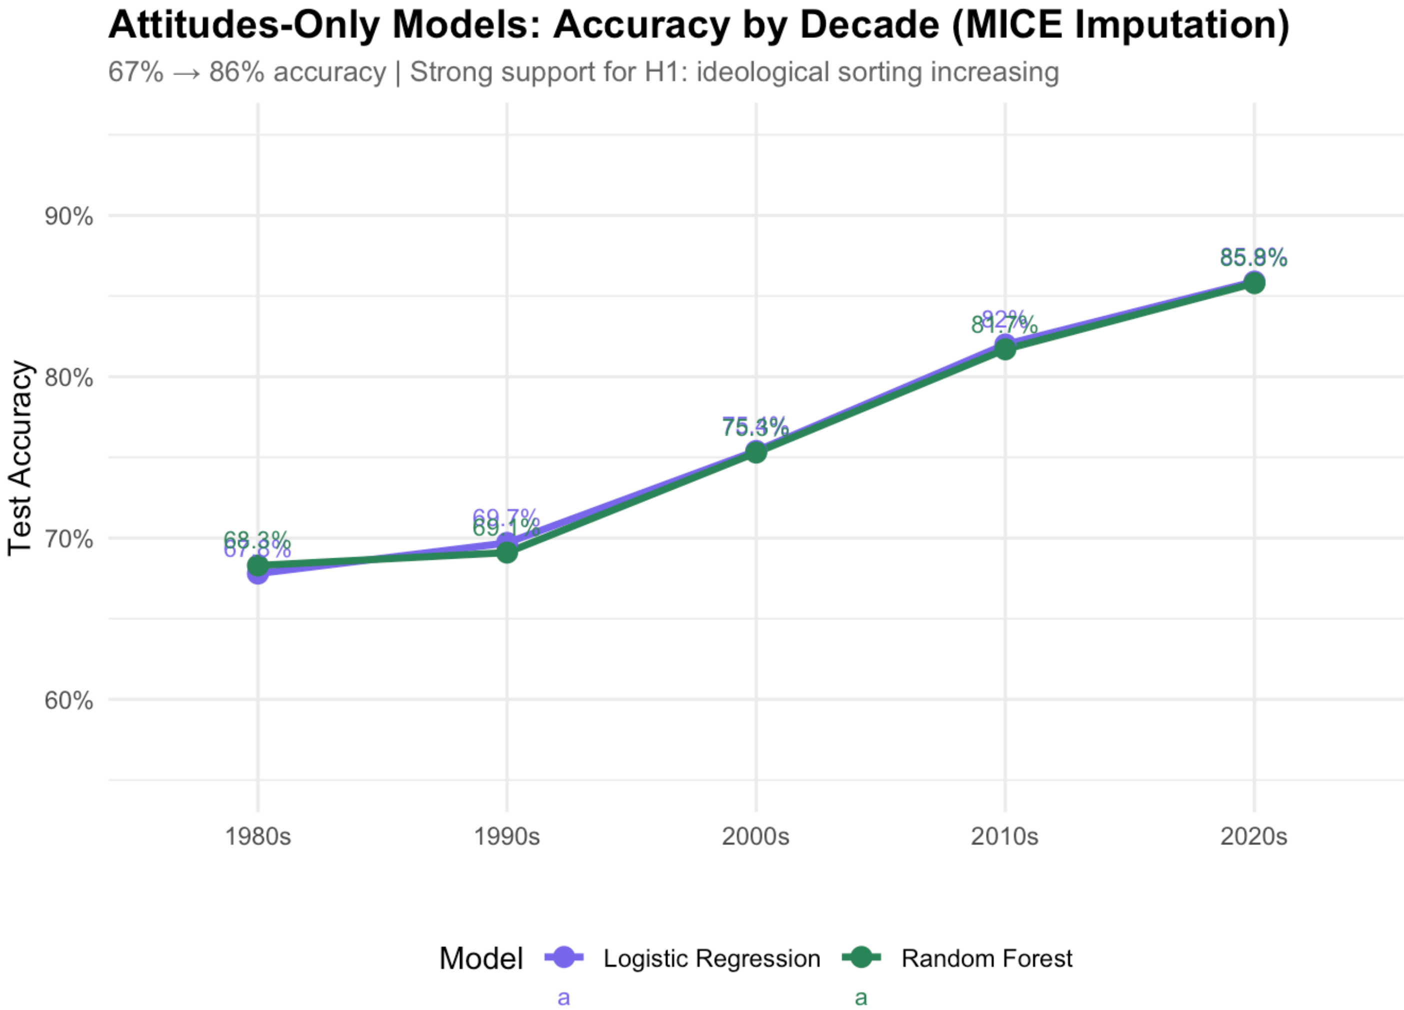

ROC Curves: 1980s vs. 2020s

1980s

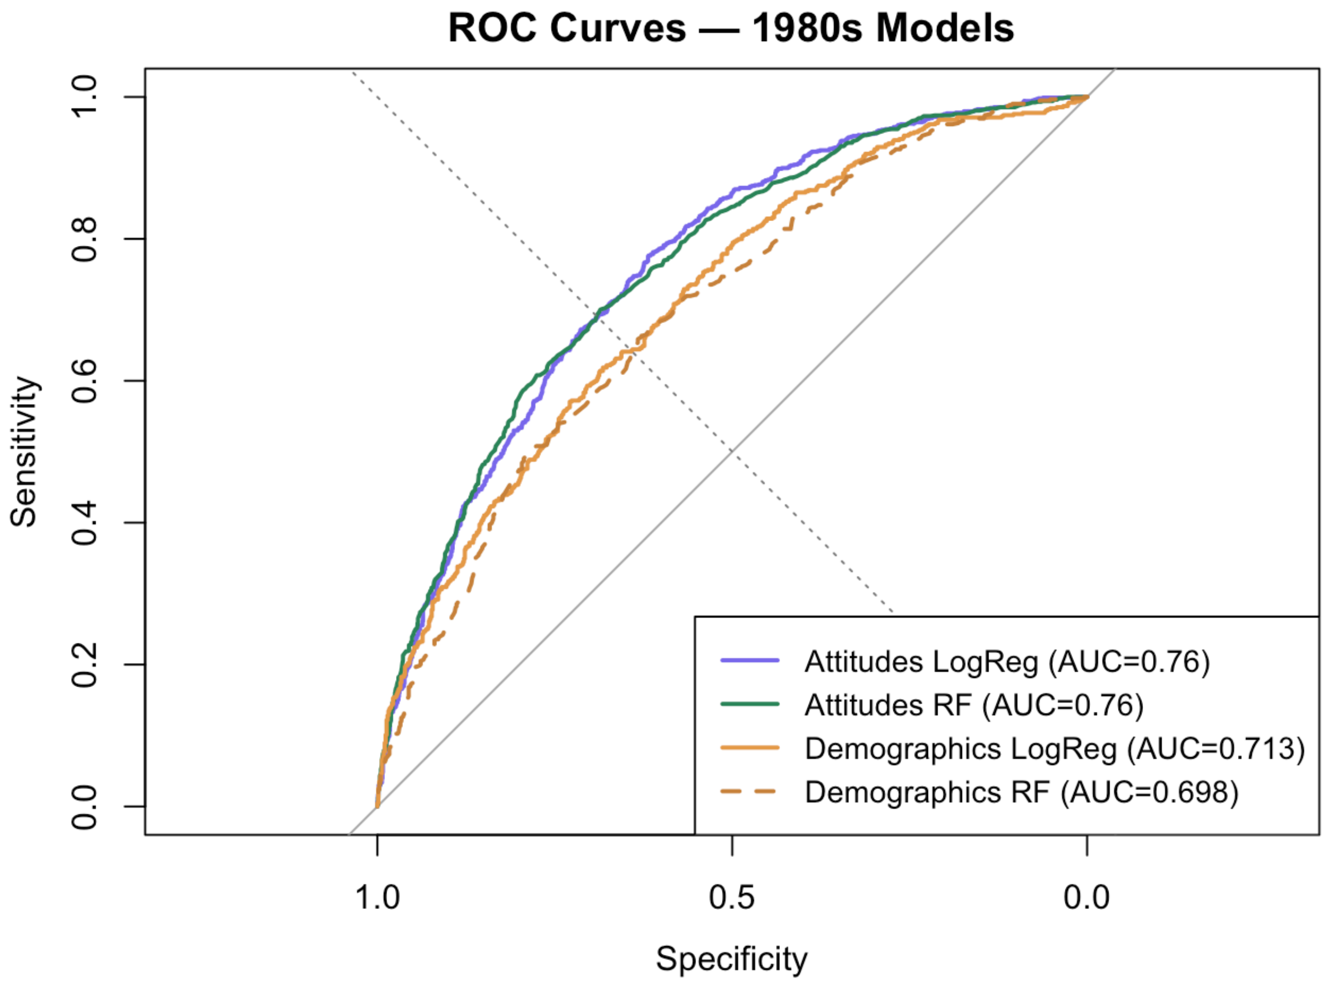
ROC Curves — 1980s Models
Sensitivity
Specificity
1.0
0.8
0.6
0.4
0.2
0.0
1.0
0.5
0.0
Attitudes LogReg (AUC=0.76)
Attitudes RF (AUC=0.76)
Demographics LogReg (AUC=0.713)
Demographics RF (AUC=0.698)

2020s

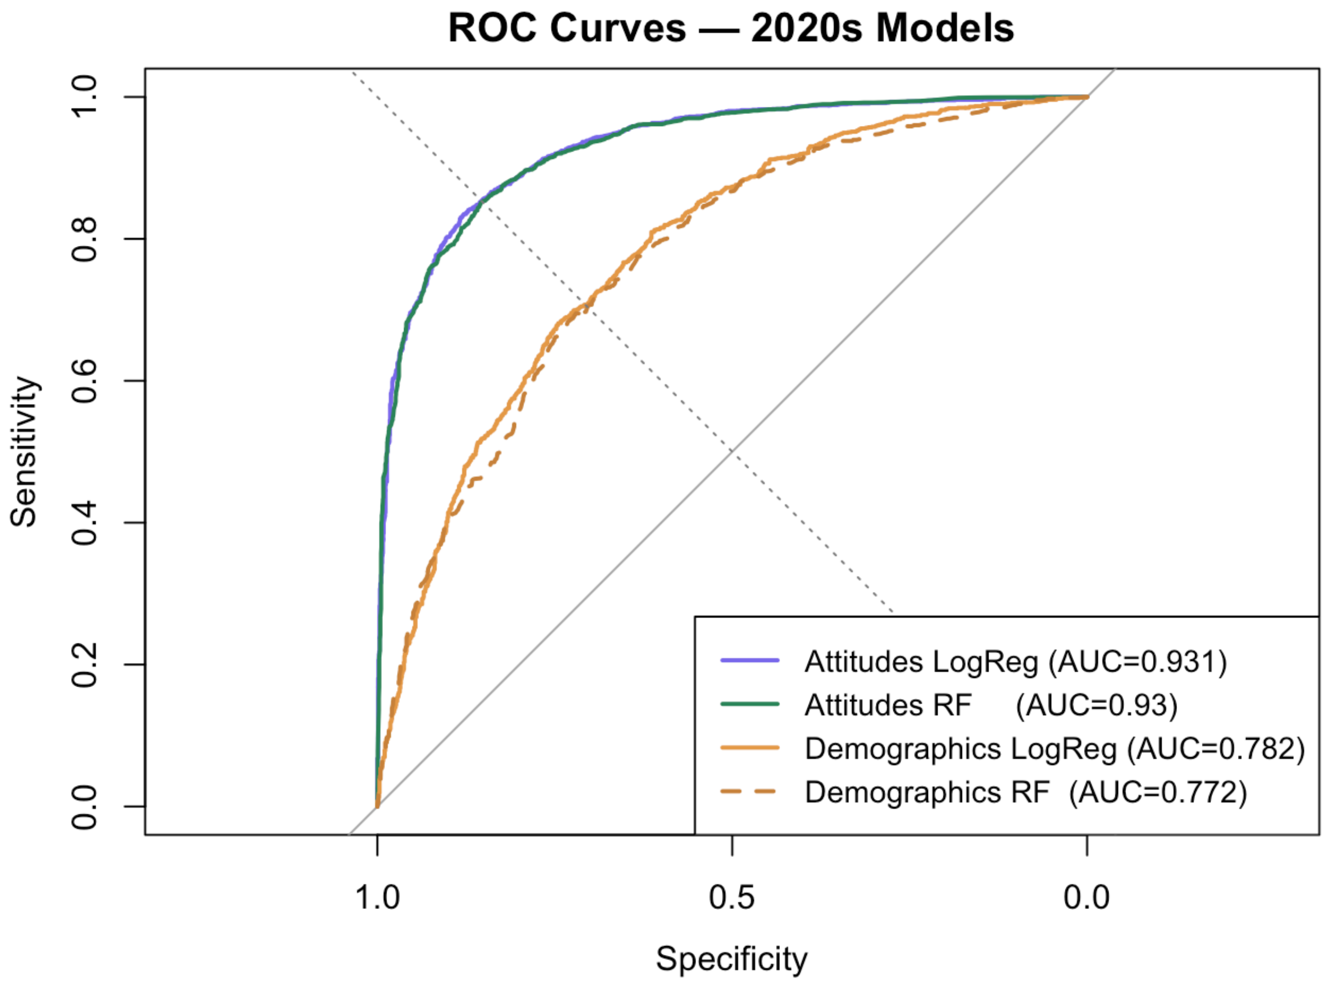

The two ROC panels show the main finding of the study. In the 1980s, the performance of the attitude-based models and demographic models were similar, with AUC values between 0.70 and 0.76. Neither was clearly better. However, by the 2020s, the situation changed significantly. The attitude-based curves (purple and green) improved to an AUC of about 0.93, while the demographic curves (orange) stayed around 0.77. The small gap from the 1980s has turned into a significant difference in the 2020s. Although demographics like race, education, and income still have some predictive power for party identification, attitudes about issues such as abortion and government spending are now much more important. Therefore, as was expected based on the reviewed literature, demographics would not become more predictive over time; instead, attitudes have become the main factor in American partisanship.

Grand Combined Plot

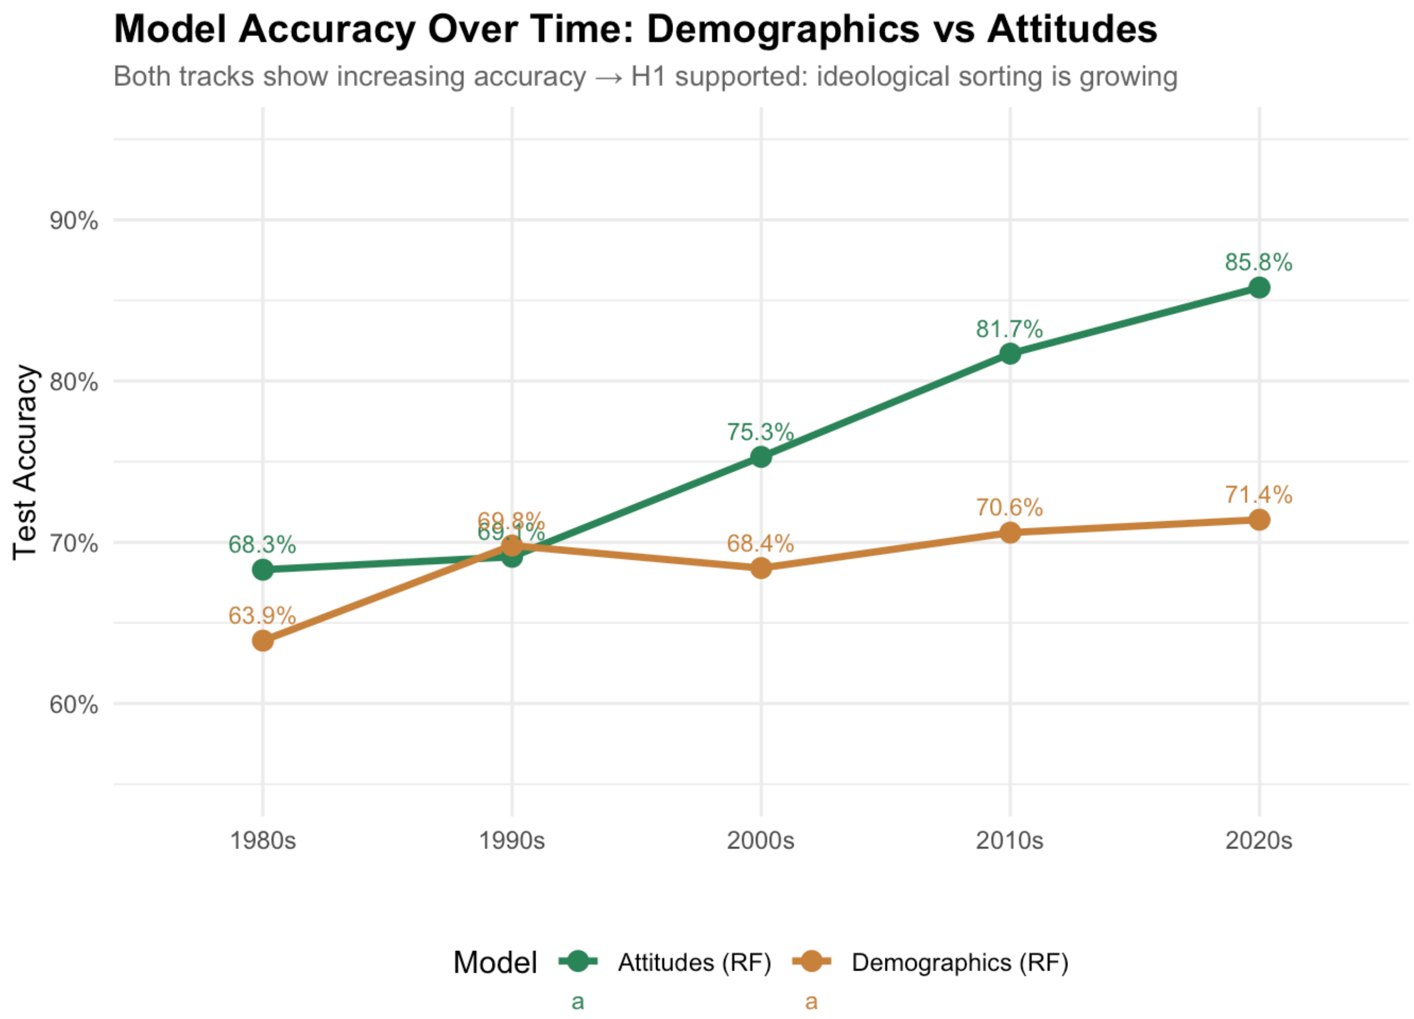

Discussion: Part 1

This analysis finds that party identification in the United States has become significantly more predictable based on people's policy views over the last four decades. In the 1980s, a model using eleven issues correctly identified about 68% of people's party affiliations; by the 2020s, this number increased to 86%. This change shows that people's beliefs about issues like abortion, government spending, and racial equality now reveal much more about their political party than they did in the past.

Demographic factors (like race, religion, education, and income) have also become slightly better at predicting party affiliation, improving from 64% to 71%. However, these factors alone still misclassify nearly one in three people. The key takeaway is that within demographic groups, individuals' attitudes and party affiliations are now much more aligned than they were 40 years ago.

Abortion stands out as the issue that has changed the most, it was not a strong predictor of party affiliation in the 1980s but became the strongest predictor by the 2020s. In contrast, economic attitudes have remained stable over the decades and haven't driven this sorting trend. This is consistent with Murphy and Zhirnov (2020), who documented the accelerating importance of cultural and moral issues in partisan alignment, particularly from the 1990s onward.

Economic attitudes have stayed mostly the same as predictors over the last five decades. While people still care about issues like government job guarantees, spending, and equality, these factors haven't driven the trend of political sorting. The changes are mainly cultural and moral, not economic. This is consistent with Murphy and Zhirnov (2020), which noted that partisan alignment has increased more in cultural issues than in other policy areas.

It is worth to note that these findings do and do not show. The models do not tell us that Americans have become more ideologically extreme, the EDA suggests that average policy positions have shifted modestly, not dramatically. What has changed is the consistency between positions and party labels. Americans have not necessarily moved further apart in what they believe; they have become more consistent in aligning their beliefs with the party that represents them. This distinction, between opinion polarization and partisan sorting, is central to the debate in the literature (Fiorina, Abrams, and Pope 2005; Abramowitz and Saunders 2008), and our findings come down clearly on the side of sorting rather than extremism as the dominant process.

Finally, the affective polarization data add an important dimension to this picture. Even as policy attitudes have become more consistently sorted, feeling thermometer scores toward the opposing party have declined steadily since the late 1970s. Americans are not just more ideologically consistent, they are also more emotionally hostile toward the other side. It's still unclear whether this emotional division is a is a cause or a consequence of ideological sorting, but both trends show a significant divide in the electorate today.

Hypothesis Evaluation

H1 — SUPPORTED. Attitude-based models have improved accuracy significantly over the decades, from about 68% in the 1980s to 86% in the 2020s. This 18-point increase suggests that people's political views are increasingly aligning with their party identities, contradicting the Fiorina, Abrams, and Pope (2005)'s argument that polarization is merely elite-driven; it is reflected in the predictability of mass partisanship from issue positions.

H2 — NOT SUPPORTED. Demographic models show only a slight improvement, rising from 64% to 71% accuracy, mainly leveling off after the 2000s. Even today, these models misclassify about 29% of respondents. This is consistent with Kim and Zilinsky (2022)'s finding that demographic characteristics reveal surprisingly little about how people vote — knowing someone's age, race, income, or education simply does not tell you much about their party. The small gains are mostly tied to religious factors, suggesting that these demographics reflect identity more than traditional voting patterns.

H3 — SUPPORTED. The data shows that issues like abortion and traditional values have become much more significant in predicting party affiliation over the years. In contrast, economic factors have remained stable and less influential, indicating that cultural and moral issues are now more important than economic ones in sorting political preferences.

Part 2: The 2024 Electorate

The first part of this analysis showed that strong partisanship has increased significantly over the last 40 years. This leads to a new question: What about those who don't identify with either major party? Who are the Democrats, Republicans, and Independents in 2024? Do Independents represent a true middle ground in attitudes, or are they just partisans without a label?

The second part of this study uses data from the 2024 ANES to explore these questions. Instead of looking at changes over time, we create a clear picture of the current voters, examining the differences in policy beliefs among the three groups and whether any patterns emerge that align with party labels without using party identification as a factor.

Data and Variable Selection

For this analysis, we are using 5,483 respondents from the 2024 ANES wave who answered valid party identification questions. This time, we include Independents as a separate group, which also covers those who lean towards a party, rather than only pure Independents. This decision is both practical and theoretical.From the practical side, there are too few pure Independents (only 380) to create meaningful comparisons. Research shows that leaners really do occupy a middle groundeven if they occasionally vote partisan (Klar and Krupnikov 2016), so they are included. Since our central question concerns attitudinal rather than behavioral independence, including leaners is the more appropriate choice. This gives us a sample of 1,930 Democrats, 1,810 Independents, and 1,743 Republicans, providing a balanced distribution.

We also expanded the attitudes we're looking at compared to Part 1. In that part, we only had 11 attitude items due to consistency issues over the decades. In the 2024 wave, we could select 17 variables with low missingness rates (below 15%) to give a better overall view of the current American political landscape as it stands today.

The choice of variables was intentional. We included one strong representative item from each significant area of American politics today. These areas cover economic attitudes (like government spending), evaluations of the political situation (like presidential approval), social issues (like abortion and immigration), racial attitudes (like aid to Black Americans), religious values, environmental and security policies, criminal justice (like the death penalty), and political engagement (like interest in elections). This way, we capture a wide range of issues that can differentiate the three partisan groups, including not just cultural issues, but also economic and governance-related attitudes.

Random Forest Classification

We begin with a random forest model trained on all 17 attitude variables and 8 demographic predictors. The model achieves an overall accuracy of 65.4% on the held-out test set, a reasonable result for a three-class classification problem, and one that already tells us something important - predicting which of three groups someone belongs to is substantially harder than the binary Democrat/Republican task from Part 1.

Random Forest Confusion Matrix (Accuracy: 65.4%)

	Democrat	Independent	Republican
Democrat	459	175	29
Independent	94	203	80
Republican	26	165	413

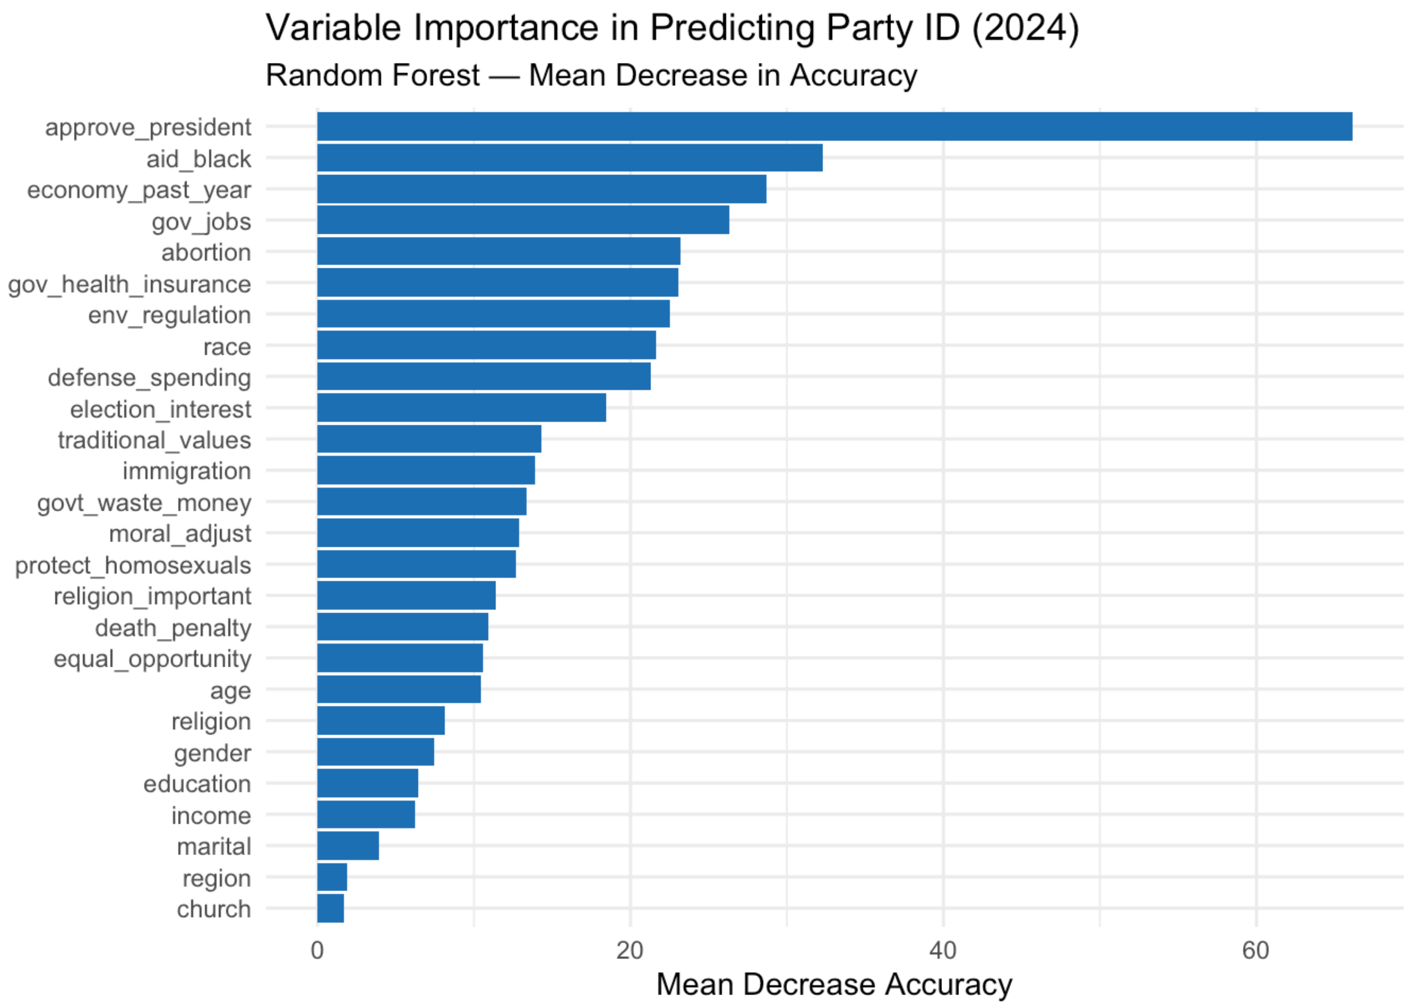

The variable importance plot shows which factors are most effective in distinguishing between the three groups. **Presidential approval stands out as the most significant factor** for 2024; how people view the current president is a better indicator of their political identity than most policy views. This aligns with the trend of American political identity becoming more centered around the president (Abramowitz and Saunders 2008). Besides presidential approval, opinions on aid to Black Americans, economic evaluations, and abortion are also important predictors. Demographic factors like race, income, and education rank lower, highlighting that attitudes about issues are more useful than demographics in predicting party identification.

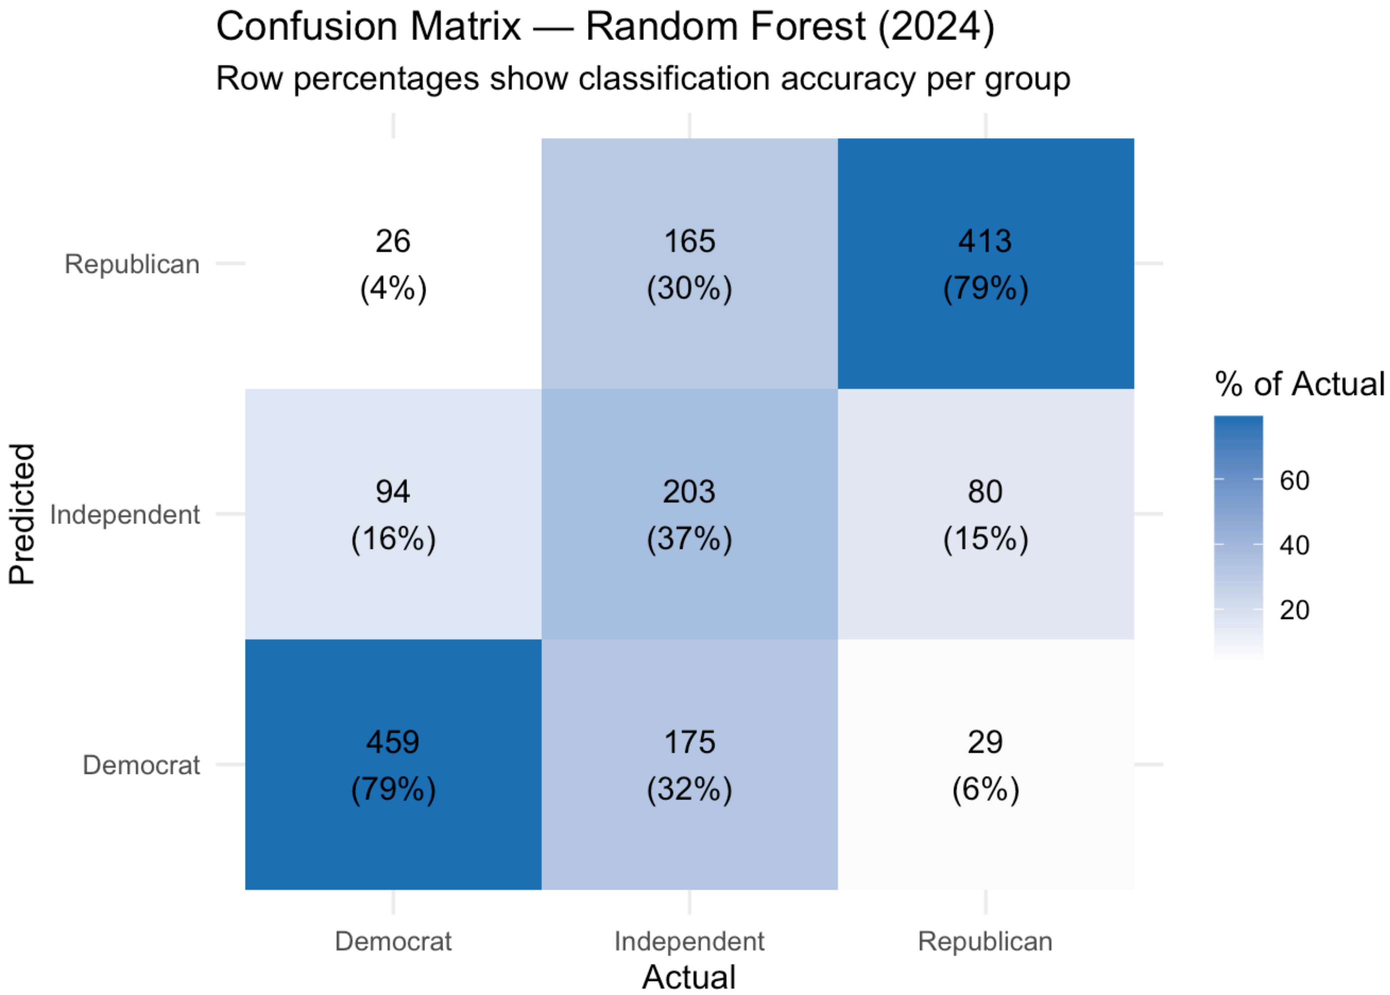

The confusion matrix shows how well the model works. It accurately classifies Democrats and Republicans, correctly identifying 459 out of 579 Democrats (79.3%) and 413 out of 522 Republicans (79.1%). However, it struggles with Independents. Out of 543 Independents, it only correctly classifies 203 (37.4%), mislabeling 340 of them as either Democrat (175) or Republican (165). This isn't a flaw in the model, it highlights that while the two parties are becoming easier to distinguish, Independents don't fit neatly into either group.

Multinomial Logistic Regression and Attitudinal Profiles

To support the random forest findings, we also used a multinomial logistic regression model, which achieved a similar accuracy of 64.2%. The comparable results from these two different models indicate that our findings are reliable and not just due to the specific method used.

Multinomial Logistic Regression Confusion Matrix (Accuracy: 64.2%)

	Democrat	Independent	Republican
Democrat	441	163	28
Independent	107	221	100

Republican	31	159	394

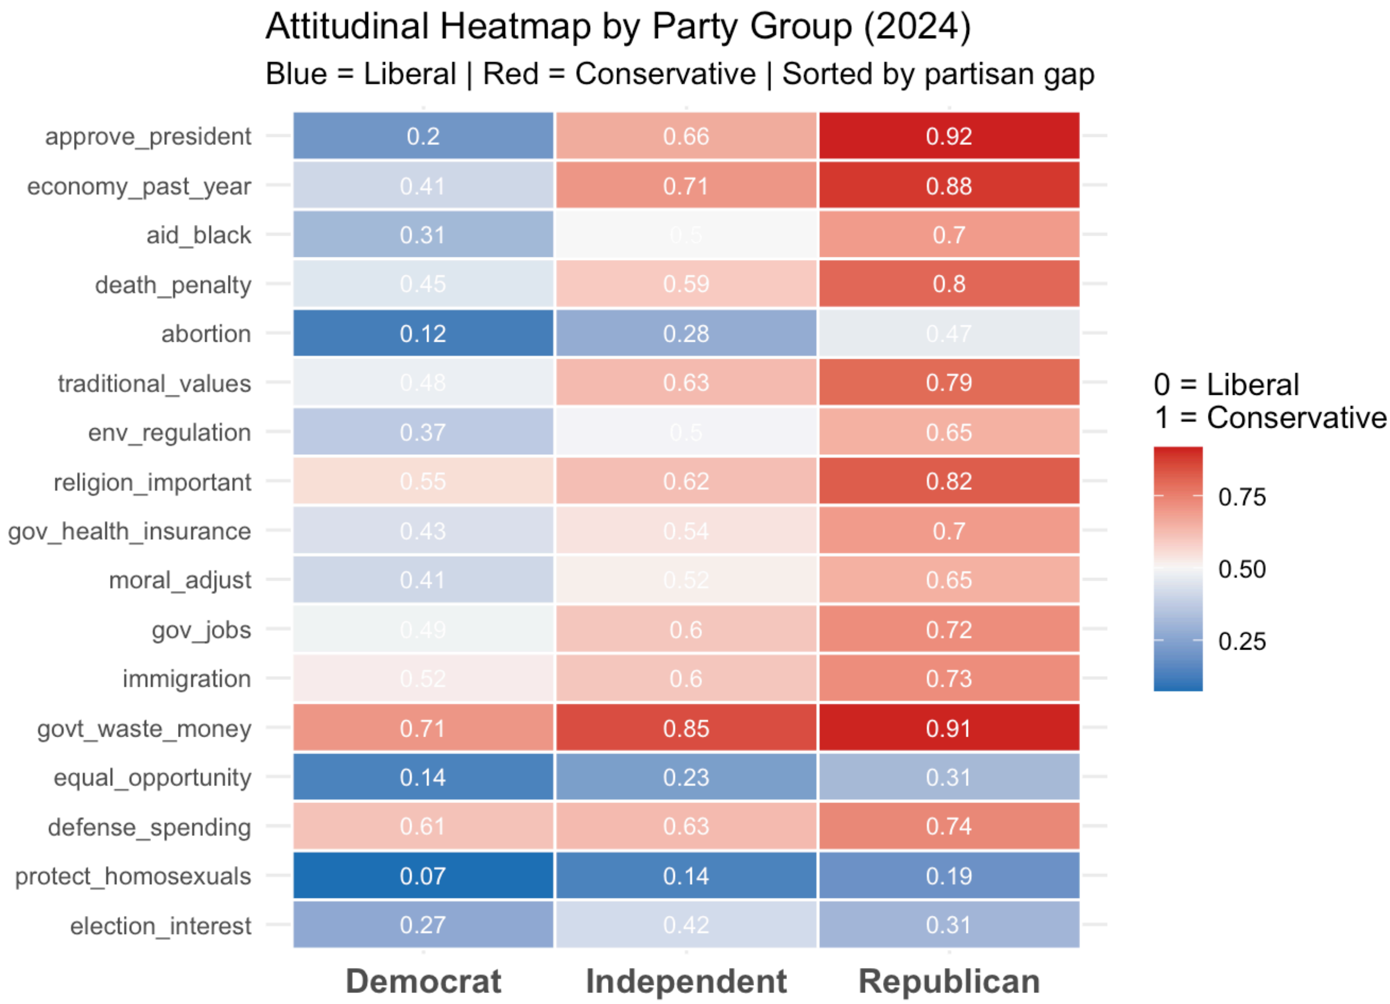

The multinomial model provides a detailed output in the form of an attitudinal heatmap, which illustrates the mean normalized position of each group on 17 attitude variables. These variables are rescaled so that a value of 0 indicates the most liberal response, while a value of 1 indicates the most conservative. The variables are arranged from top to bottom based on the size of the gap between Democrats and Republicans, with the issues that create the largest divisions between the two parties appearing at the top.

Several things are interesting. Democrats and Republicans have very different views on presidential approval and economic evaluations, showing big partisan gaps. This indicates that people's opinions about the president are closely tied to their political identity today. Other issues like abortion, support for Black Americans, and environmental regulations also show consistent differences, with Democrats generally holding more liberal views and Republicans leaning more conservative.

In addition, traditional values, religion, and immigration follow a similar trend, with Republicans scoring significantly higher on all three topics. This aligns with the culture war sorting discussed in Part 1 and supported by existing literature (Baldassarri and Gelman 2008; Murphy and Zhirnov 2020).

The column about Independents is the most revealing. They often fall between Democrats and Republicans on various issues, showing that there is a true middle ground in the 2024 electorate. However, their views differ greatly depending on the topic. For example, Independents are more critical of the Biden administration and lean toward Republican views on presidential approval and the economy. On social issues like gay rights and equal opportunity, they are closer to Democrats. This shows that Independents have a mixed profile, being more conservative on some issues and more liberal on others.

Misclassification Analysis

The confusion matrix already established that Independents are misclassified at a far higher rate than partisans — 62.6% compared to roughly 21% for both Democrats and Republicans. But the more interesting question is not that the model struggles with Independents, but who exactly is being misclassified and why.

We looked at how well people's political views match their party affiliation by examining various attitudes and demographics. We found five key factors (statistically significant) that affect this. First, people who are less interested in elections tend to be misclassified as having inconsistent party views. Second, women are less likely to be misclassified compared to men, meaning their views are more closely aligned with their party. Third, race plays a significant role, showing changing relationships between race and party alignment. Marital status has a small effect, and views on abortion are important; those with unclear positions on abortion are harder to classify, reflecting its significant role in political divisions for 2024.

Together these findings suggest that misclassification is not random. The people the model struggles with **tend to be less politically engaged, hold more ambiguous views on cultural issues, and come from demographic groups that are changing their political loyalties**. This reflects real uncertainty in their attitudes, rather than just being moderate.

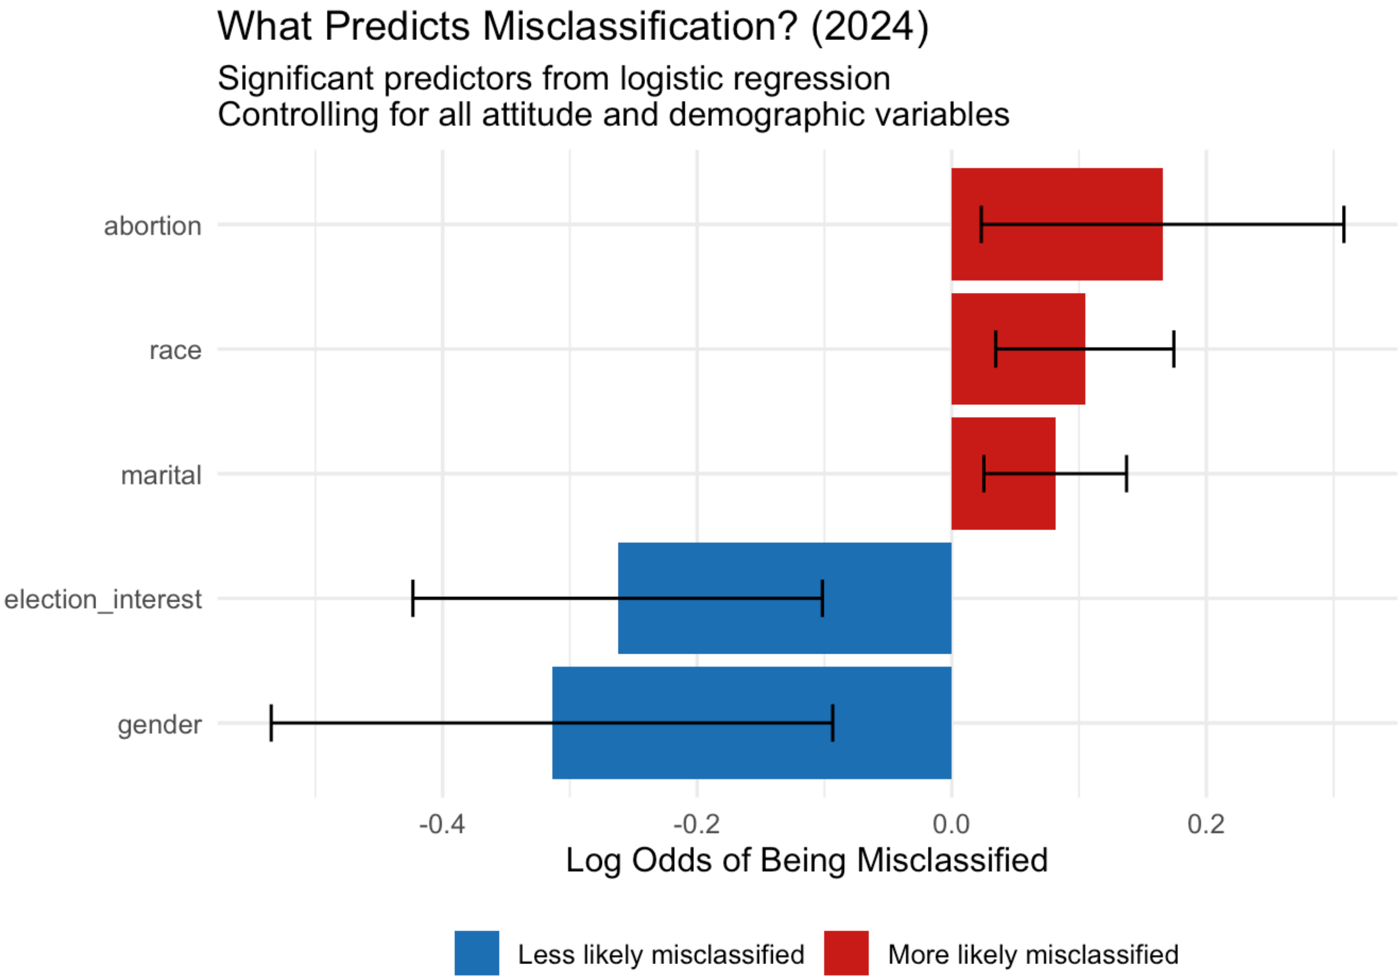

Latent Class Analysis

In our previous analyses, we focused on party identification as the main outcome. However, in this Latent Class Analysis, we take a different approach by using only attitude variables without any party labels. This allows the model to identify natural groupings on its own. If the model performs well, the groups it identifies should align closely with Democrats, Independents, and Republicans, even though it has no prior knowledge of these categories.

We tested solutions with two, three, and four classes. The three-class solution was the best fit, as the BIC significantly decreased from the two-class model, while the four-class model produced a small and difficult-to-interpret group. Therefore, we decided to proceed with three classes.

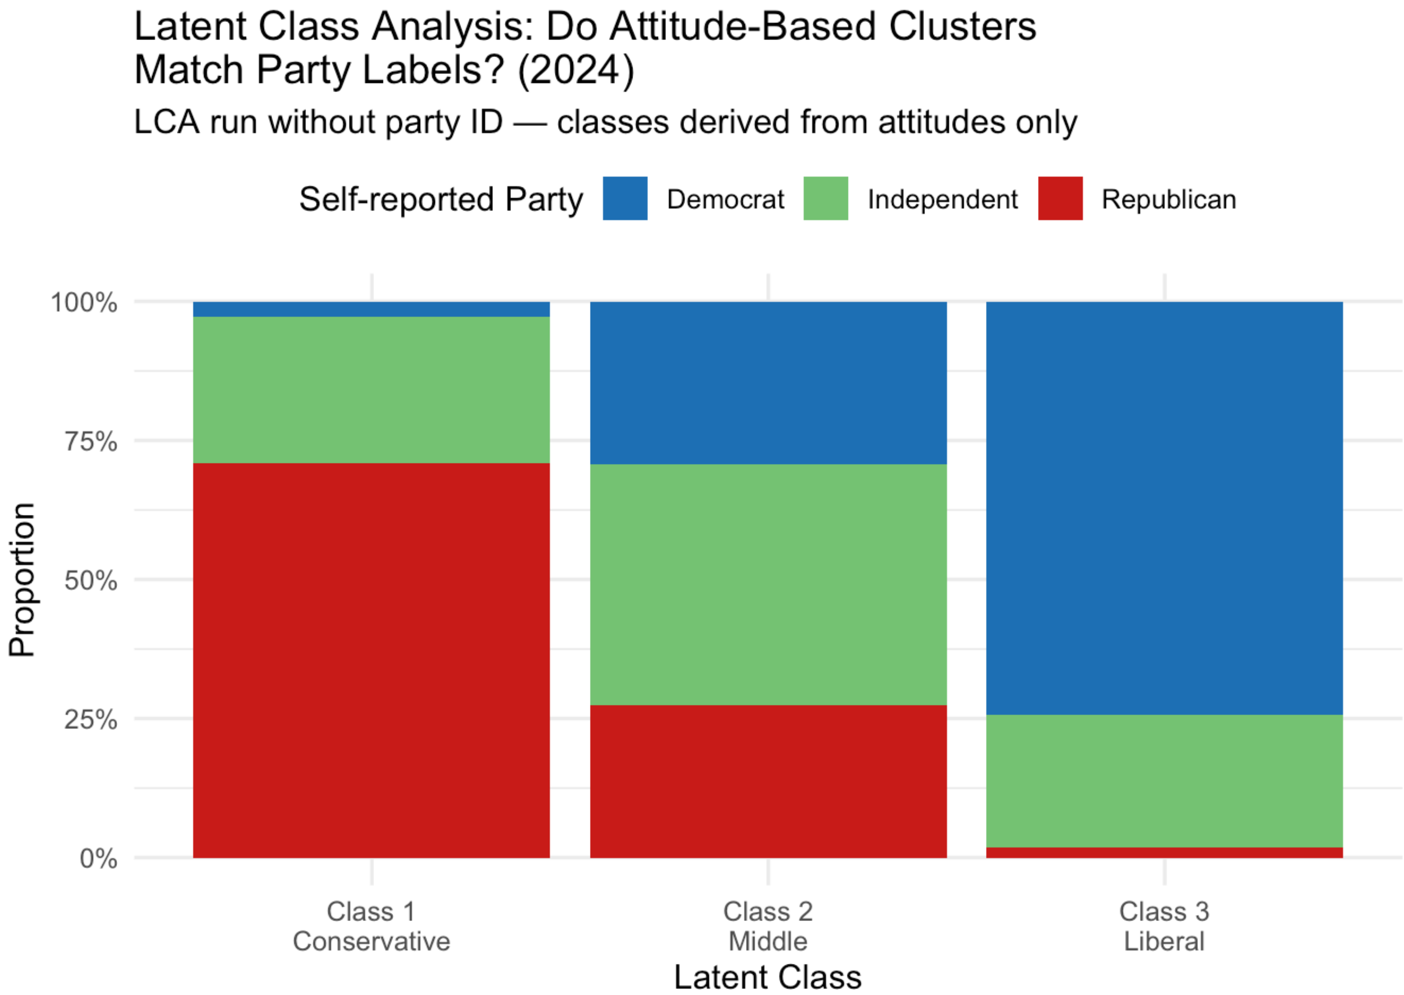

The results are very interesting. Class 1 (27.6% of respondents) is 70.9% Republican. Class 3 (30.4%) is 74.4% Democrat. The model recovers something very close to the partisan divide using attitudes alone, without ever seeing party identification. Class 2 (42%) is different, it is 43.5% Independent, with Democrats and Republicans split roughly equally for the rest. This is the attitudinal middle. But the most important finding is where Independents end up. They are spread across all three classes, not concentrated in the middle. Some Independents hold clearly conservative attitudes and land in Class 1. Others hold clearly liberal attitudes and land in Class 3. Only the largest share ends up in the ambiguous middle class. This directly addresses H4: **Independents are not a single coherent group. Some are effectively partisans who reject the label. Others genuinely sit between the two parties in their views.**

Demographic Portrait of Attitudinal Classes

To understand who belongs to each attitude group, we analyze the demographic makeup of every group.

The three groups have some clear differences. The Liberal group is the youngest, with an average age of 52.1 years, and they are the most educated, about 64.3% have a college degree, while only around 37-40% of the others do. Members of this group are spread out fairly evenly across different regions, although there are fewer Liberals in the South. In terms of religion, nearly 60% identify as Other or None, which means they are the least religious of all three groups.

The Conservative group is the oldest, averaging 58.0 years. They tend to be more religious, with 44.7% identifying as Catholic and 18.5% as Protestant. This group is mostly located in the South, where 44.9% of its members live, and they have the smallest percentage of college graduates.

The Middle group falls in between the two others in most areas, with an average age of 50.9 years and about 36.7% being college educated. Their religious makeup is also somewhere in the middle. Notably, they have the highest percentage of Catholics among the three groups and are largely found in the South and Midwest.

The demographic differences show distinct attitudes. The Liberal group tends to align with those who support the Democratic Party, often being more educated and secular. The Conservative group is associated with the Republican Party, consisting of older, more religious individuals, mainly from the South. The Middle group is more varied and difficult to categorize.

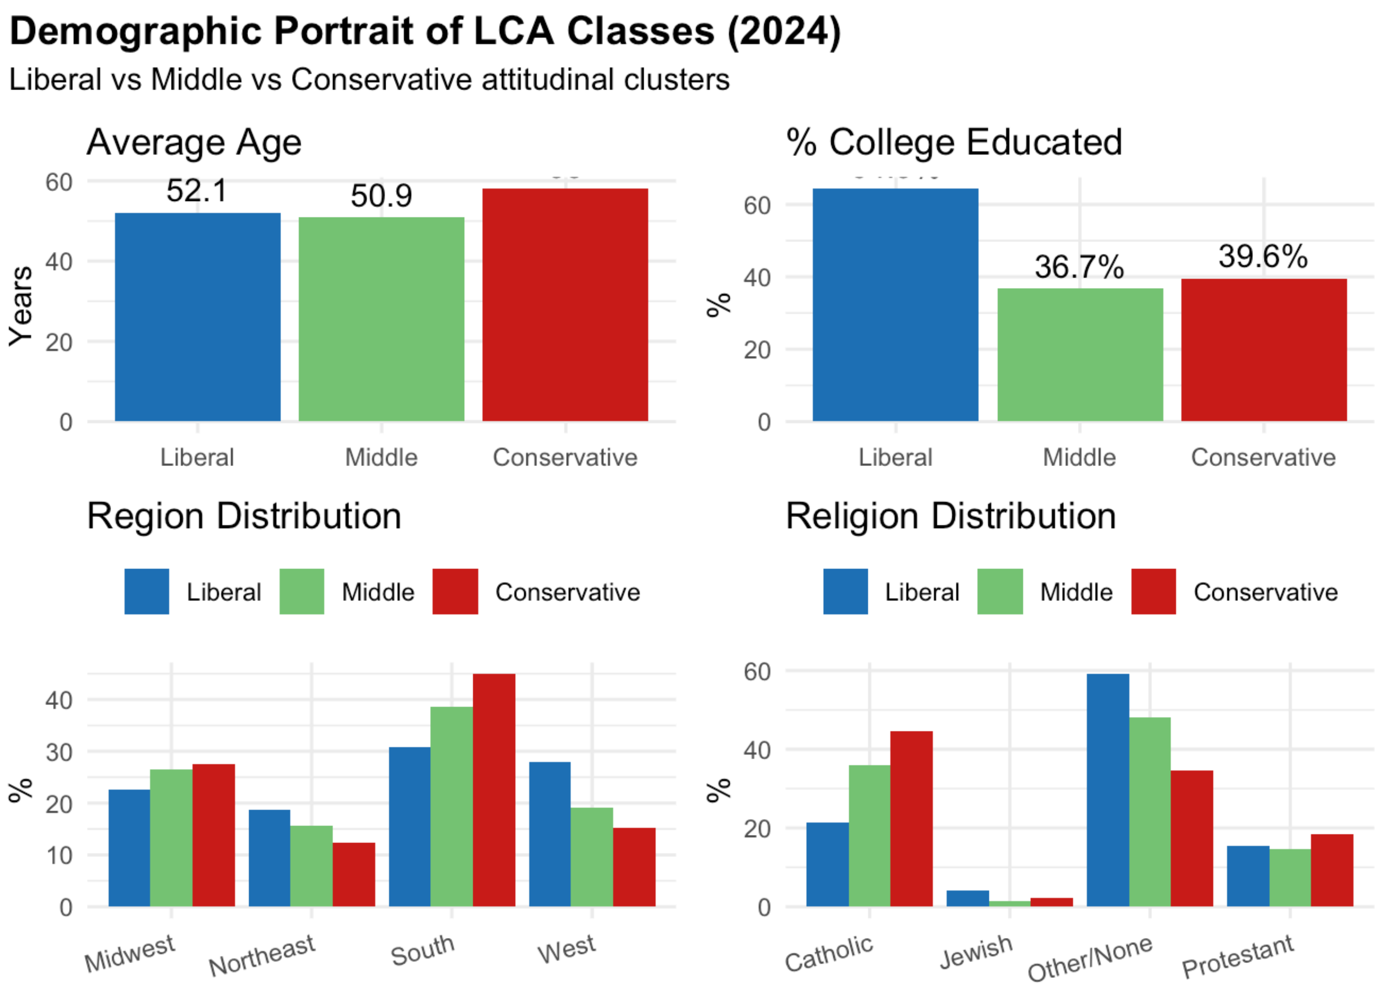

Discussion: Part 2

The 2024 analysis builds on the findings from Part 1. If we see that strong supporters of both parties are clearly defined, the next question is: what about everyone else?

The answer is that Democrats and Republicans can be predicted based on their opinions—about 80% of the time, we can correctly guess their political leanings. However, it's a different story for Independents. The model misclassifies 62.6% of them, and this isn't due to any technical issues. It shows that Independents often don't have a consistent set of beliefs that make them different from partisans.

The attitudinal heatmap illustrates this. On many issues, Independents tend to fall between Democrats and Republicans, but the gap can vary. For example, on presidential approval and views about the economy, they lean more toward Republicans. On social issues, like gay rights, they lean more toward Democrats. This aligns with the idea that rather than being steady in the middle, many Independents favor one party or the other based on the topic.

The Latent Class Analysis supports this idea. When we group people based strictly on their opinions, without considering party labels, we see that Independents don't form a single group. The biggest share ends up in a mixed category, but many also fit into clearly conservative or clearly liberal groups. The Independent label includes at least two types of people: some truly sit between the two parties, while others have strong partisan views but prefer not to identify as either Democrat or Republican.

H4 — PARTIALLY SUPPORTED. The data confirm that most Independents do not occupy a stable middle ground, their views tilt toward one party or the other depending on the issue. However, H4 also predicted that most Independents would align more closely with one of the two parties than with a genuine middle. The LCA complicates this, while many Independents do lean partisan attitudinally, the single largest group lands in the ambiguous middle class rather than in either partisan cluster. The more accurate conclusion is that Independents are not one group at all, some are effectively undercover partisans, while others genuinely sit between the two parties in ways that neither party fully captures.

H5 - SUPPORTED. Those who were hardest to classify tended to show lower levels of political engagemen. This aligns with Klar and Krupnikov's point that many Independents choose not to identify with a party not because they are moderate, but because they prefer to stay away from political conflict.

Conclusion

This study explores two main questions: How have people's attitudes and their political party affiliation changed since the 1980s? And what does the makeup of the 2024 voters look like?

For the first question, the answer is quite clear. People have become more sorted in their political beliefs. Today, attitude-based models can predict someone's party affiliation with 86% accuracy, compared to just 68% four decades ago. This improvement is more than twice what we see with demographic models over the same time frame. In fact, knowing someone's age, income, education, or race helps us only about two-thirds of the time, meaning one in three people might still be wrongly classified. What truly predicts party affiliation is not who you are, but what you believe. While Americans haven't necessarily become more extreme, their views and party labels have become much more closely aligned. Issues related to culture, especially abortion, have become stronger indicators of party affiliation, overshadowing economic issues.

On the second question whether a genuine middle ground exist the answer is partly. Among strong partisans, sorting is nearly complete. Democrats and Republicans show clear and recognizable attitudinal profiles, with the largest gaps on presidential approval, economic evaluations, abortion, and immigration. Independents don't fit neatly into one category because their views are a mix of both Democratic and Republican opinions. When we look at their attitudes without party labels, we see that they don't form a single group. Some are clearly conservative, others are clearly liberal, and many are in the middle. While there is a true middle group, it only represents part of the Independent voters. The rest are more like partisans who just don’t want to identify with a party.

Using demographic and opinion data, we created a broad picture of today’s Independents and it shows that they are not just a single political group but a mix of people who don’t completely align with either major party.In summary, while the electorate in 2024 is more sorted among dedicated party members, it is more ambiguous for those who identify as Independents. Understanding who Independents really are and whether they represent a steady political identity or something more fluid is a crucial question in American politics.

References

Abramowitz, Alan I., and Kyle L. Saunders. 2008. “Is Polarization a Myth?” *Journal of Politics* 70 (2): 542–55. https://doi.org/10.1017/S0022381608080493 (https://doi.org/10.1017/S0022381608080493).

Baldassarri, Delia, and Andrew Gelman. 2008. “Partisans Without Constraint: Political Polarization and Trends in American Public Opinion.” *American Journal of Sociology* 114 (2): 408–46. https://doi.org/10.1086/590649 (https://doi.org/10.1086/590649).

Fiorina, Morris P., Samuel J. Abrams, and Jeremy C. Pope. 2005. *Culture War? The Myth of a Polarized America*. New York: Pearson Longman.

Gallup. 2024. “Party Affiliation.” https://news.gallup.com/poll/15370/party-affiliation.aspx (https://news.gallup.com/poll/15370/party-affiliation.aspx).

Iyengar, Shanto, Yphtach Lelkes, Matthew Levendusky, Neil Malhotra, and Sean J. Westwood. 2019. “The Origins and Consequences of Affective Polarization in the United States.” *Annual Review of Political Science* 22: 129–46. https://doi.org/10.1146/annurev-polisci-051117-073034 (https://doi.org/10.1146/annurev-polisci-051117-073034).

Kim, Seo-young Silvia, and Jan Zilinsky. 2022. “Division Does Not Imply Predictability: Demographics Continue to Reveal Little about Voting and Partisanship.” *Political Behavior* 45: 1101–26. https://doi.org/10.1007/s11109-022-09816-z (https://doi.org/10.1007/s11109-022-09816-z).

Klar, Samara, and Yanna Krupnikov. 2016. *Independent Politics: How American Disdain for Parties Leads to Political Inaction*. Cambridge University Press.

Mason, Lilliana. 2018. *Uncivil Agreement: How Politics Became Our Identity*. Chicago: University of Chicago Press.

Murphy, Justin, and Andrei Zhirnov. 2020. “Issue Partisanship and Partisan Constraint.” *Political Behavior* 42 (4): 1155–78. https://doi.org/10.1007/s11109-019-09532-z (https://doi.org/10.1007/s11109-019-09532-z).

Ojer, Maria, Mirta Galesic, Thomas T. Hills, and Alfredo J. Morales. 2025. “Divergence and Convergence: Multidimensional Ideological Polarization Across Social Groups.” *Proceedings of the National Academy of Sciences* 122 (3). https://doi.org/10.1073/pnas.2409507122 (https://doi.org/10.1073/pnas.2409507122).